

# Parallele und Verteilte Systeme

**Grundlagen zur Programmierung von Mehrkern-Systemen,  
verteilter Berechnung und massiver Parallelität**



Vorlesung WS 2017/18

(Unterlagen nur für den internen Gebrauch!)

# Übersicht

- **Massiv parallele Programmierung**
  - Optimierung der Matrix-Multiplikation
    - Mehrdimensionale Gitter (NDGRID)
  - Parallelreduktion
  - Speicherorganisation
  - Aufteilung der Kernel-Aufrufe
  - Vektordatentypen
  - Sonstige Funktionen

# Matrix-Multiplikation (zeilenweise)

$$\begin{matrix} P_1 & & & & \\ P_2 & & & & \\ P_3 & & & & \\ P_1 & & & & \\ P_2 & & & & \\ P_3 & & & & \\ P_1 & & & & \end{matrix} = \begin{matrix} P_1 & & & & \\ P_2 & & & & \\ P_3 & & & & \\ P_1 & & & & \\ P_2 & & & & \\ P_3 & & & & \\ P_1 & & & & \end{matrix} \cdot \begin{matrix} P_1 & & & & \\ P_2 & & & & \\ P_3 & & & & \\ P_1 & & & & \\ P_2 & & & & \\ P_3 & & & & \\ P_1 & & & & \end{matrix}$$

**C**                    **A**                    **B**

# Version 1: C zeilenweise

```
const char *KernelSource =
"#define DIM 1000                                     // Groesse der Matrix \n"
"__kernel void matmult(__global float *A, __global float *B, __global float *C)\n"
"{
"    int i, j, k;
"    i = get_global_id(0);
"    for (j = 0; j < DIM; j++)
"        for (k = 0; k < DIM; k++)
"            C[i*DIM+j] += A[i*DIM+k] * B[k*DIM+j];
"
"}\n";
```

# Änderungen am Hauptprogramm

```
#define DIM      1000                                // Matrix Dimensionen
#define DATA_SIZE DIM*DIM*sizeof(float)               // Matrixgröße im Speicher

void main(int argc, char **argv)
{
    // ...
    float **A, **B, **C;                            // Matrizen
    cl_mem Ap, Bp, Cp;
    size_t global[1] = {DIM};                         // Oder: size_t global = DIM;

    A = alloc_mat(DIM, DIM); init_mat(A, DIM, DIM);   // Speicher für Matrizen holen
    B = alloc_mat(DIM, DIM); init_mat(B, DIM, DIM);   // und initialisieren
    C = alloc_mat(DIM, DIM);
    // ...

    kernel = clCreateKernel(program, "matmult", &err);           // Spezifizierte Kernel
    Ap = clCreateBuffer(context, CL_MEM_READ_ONLY, DATA_SIZE, NULL, &err); // Erzeuge Puffer
    Bp = clCreateBuffer(context, CL_MEM_READ_ONLY, DATA_SIZE, NULL, &err);
    Cp = clCreateBuffer(context, CL_MEM_READ_WRITE, DATA_SIZE, NULL, &err);
    clEnqueueWriteBuffer(command_queue, Ap, CL_TRUE, 0, DATA_SIZE, A[0], 0, NULL, NULL);
    clEnqueueWriteBuffer(command_queue, Bp, CL_TRUE, 0, DATA_SIZE, B[0], 0, NULL, NULL);
    clSetKernelArg(kernel, 0, sizeof(cl_mem), &Ap); // Setzte die Argumentliste für Kernel
    clSetKernelArg(kernel, 1, sizeof(cl_mem), &Bp);
    clSetKernelArg(kernel, 2, sizeof(cl_mem), &Cp);
    // ...
    clEnqueueReadBuffer(command_queue, Cp, CL_TRUE, 0, DATA_SIZE, C[0], 0, NULL, NULL);
    // ...
}
```



# Optimierungspotential?

Matrix-Multiplikation mit OpenCL

# Version 2: C zeilenweise, schreibend

```
const char *KernelSource =  
"#define DIM 1000                                     // Groesse der Matrix \n"  
"__kernel void matmult(__global float *A, __global float *B, __global float *C) { \n"  
"    int i, j, k;  
"    float sum;                                         // Private Variable für Zwischenergebnisse \n"  
"    i = get_global_id(0);  
"    for (j = 0; j < DIM; j++) {  
"        sum = 0.0;  
"        for (k = 0; k < DIM; k++)  
"            sum += A[i*DIM+k] * B[k*DIM+j];  
"        C[i*DIM+j] = sum;  
"    }  
"}  
\n";
```

```
void main(int argc, char **argv)  
{  
// ...  
Cp = clCreateBuffer(context, CL_MEM_WRITE_ONLY, DATA_SIZE, NULL, &err);
```

# Matrix-Multiplikation (elementweise)

P <sub>1</sub>	P <sub>2</sub>	P <sub>3</sub>	P <sub>1</sub>	P <sub>2</sub>
P <sub>3</sub>	...			
			P <sub>1</sub>	P <sub>2</sub>

C

=


A


B

# Version 3: C elementweise

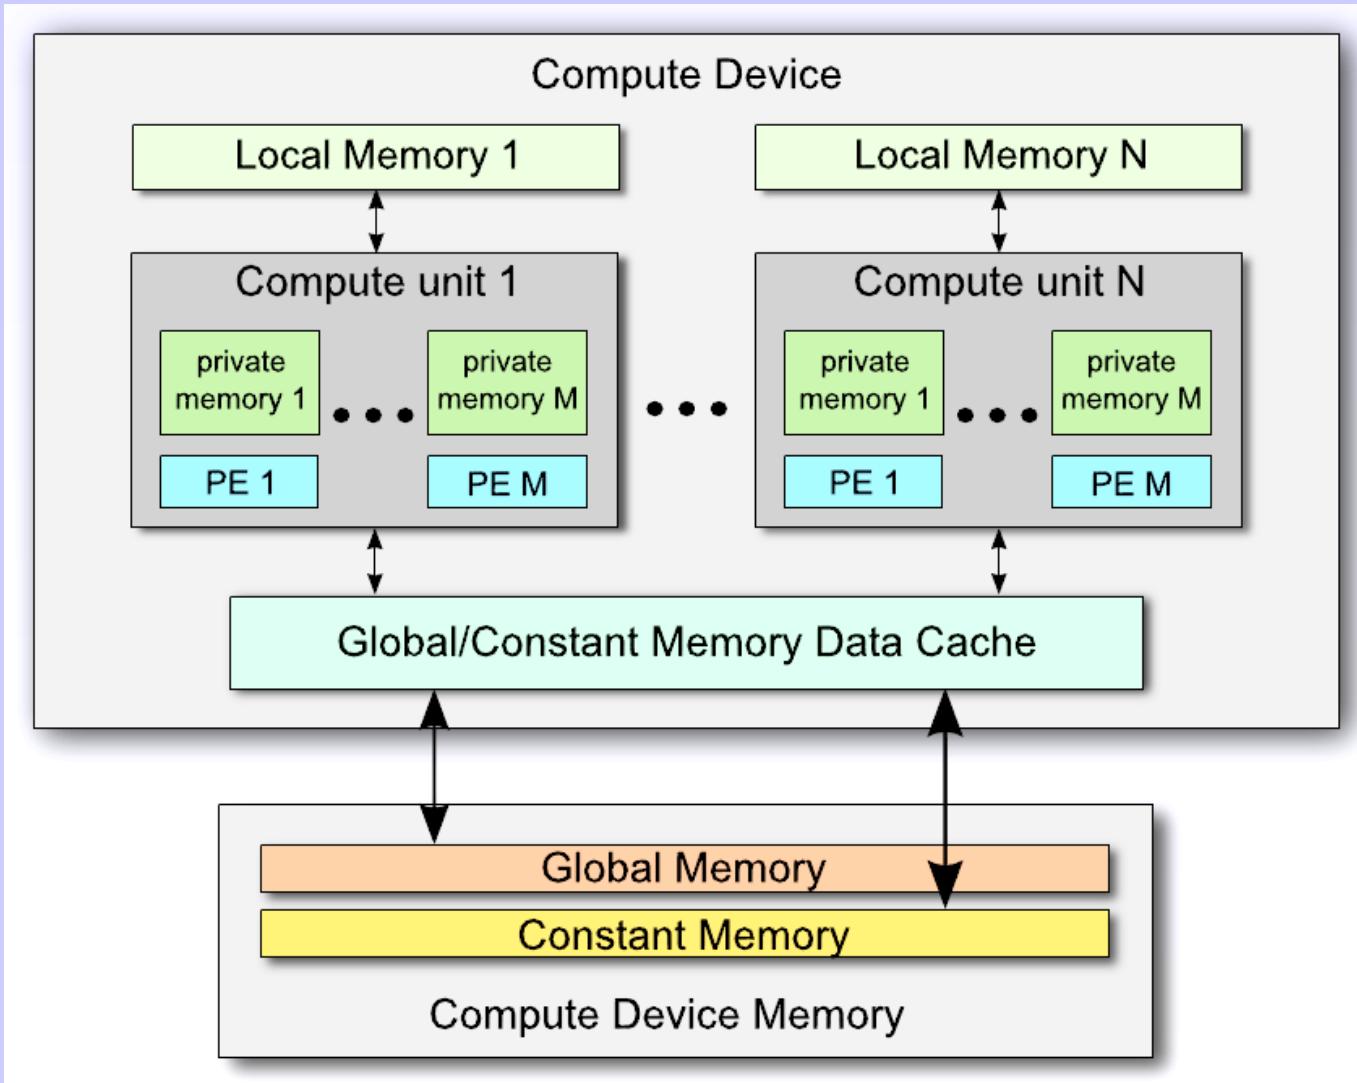
```
const char *KernelSource =
"#define DIM 1000                                     // Size of matrix \n"
"__kernel void matmult(__global float *A, __global float *B, __global float *C) { \n"
"    int i, j, k;                                     \n"
"    float sum = 0.0;                                  \n"
"    i = get_global_id(0);                            \n"
"    j = get_global_id(1);                            \n"
"    for (k = 0; k < DIM; k++)                      \n"
"        sum += A[i*DIM+k] * B[k*DIM+j];           \n"
"    C[i*DIM+j] = sum;                             \n"
"}                                              \n"
"\n";
```

```
void main(int argc, char **argv)
{
    size_t global[2] = {DIM, DIM};
    // ...
    clEnqueueNDRangeKernel(command_queue, kernel, 2, NULL, global, NULL, 0, NULL, NULL);
```

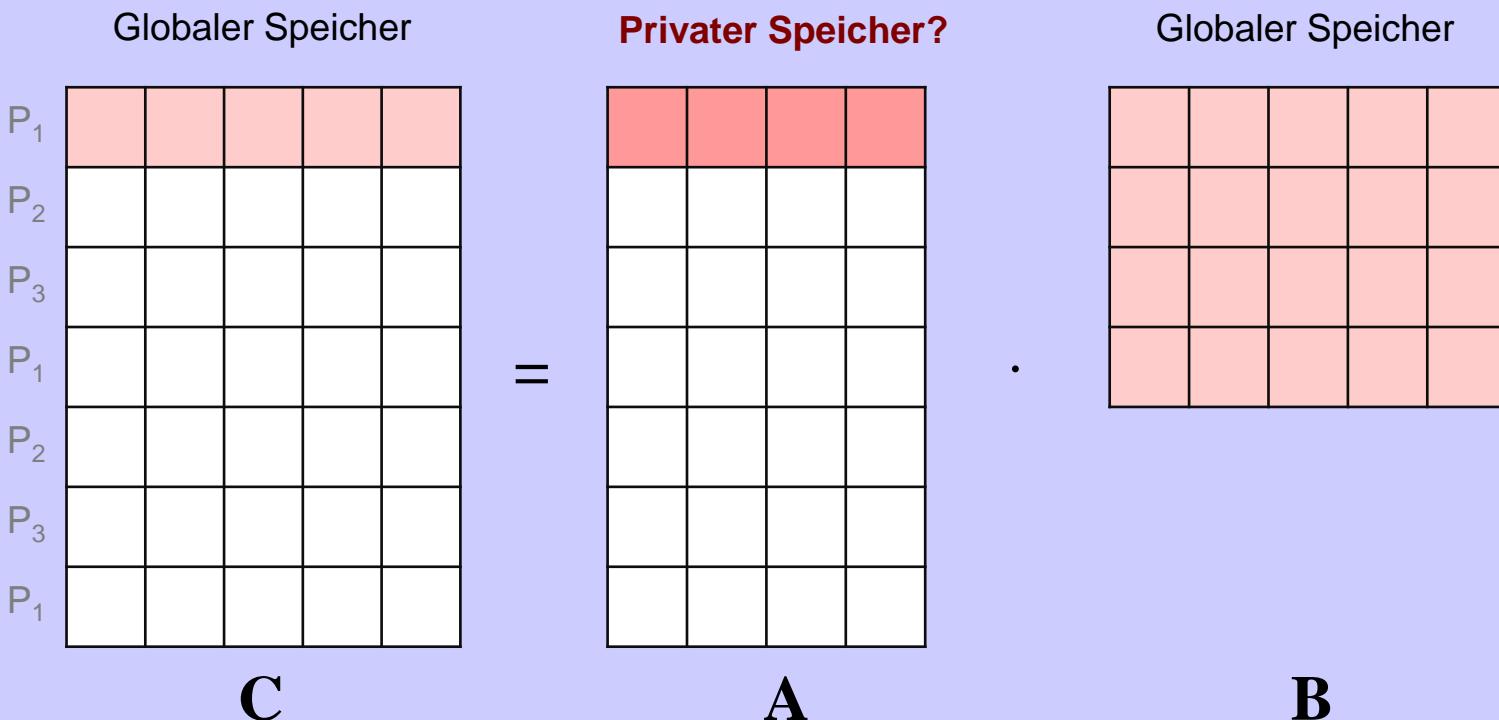
# Version 4: Schleifen tauschen?

```
const char *KernelSource =
"#define DIM 1000                                     // Size of matrix \n"
"__kernel void matmult(__global float *A, __global float *B, __global float *C) { \n"
"    int i, j, k;                                     \n"
"    float sum = 0.0;                                   \n"
"    j = get_global_id(0);                           \n"
"    i = get_global_id(1);                           \n"
"    for (k = 0; k < DIM; k++)                         \n"
"        sum += A[i*DIM+k] * B[k*DIM+j];             \n"
"    C[i*DIM+j] = sum;                                \n"
"}                                                 \n"
"\n";
```

# Speicheroptimierung?



# Matrix-Multiplikation (zeilenweise)

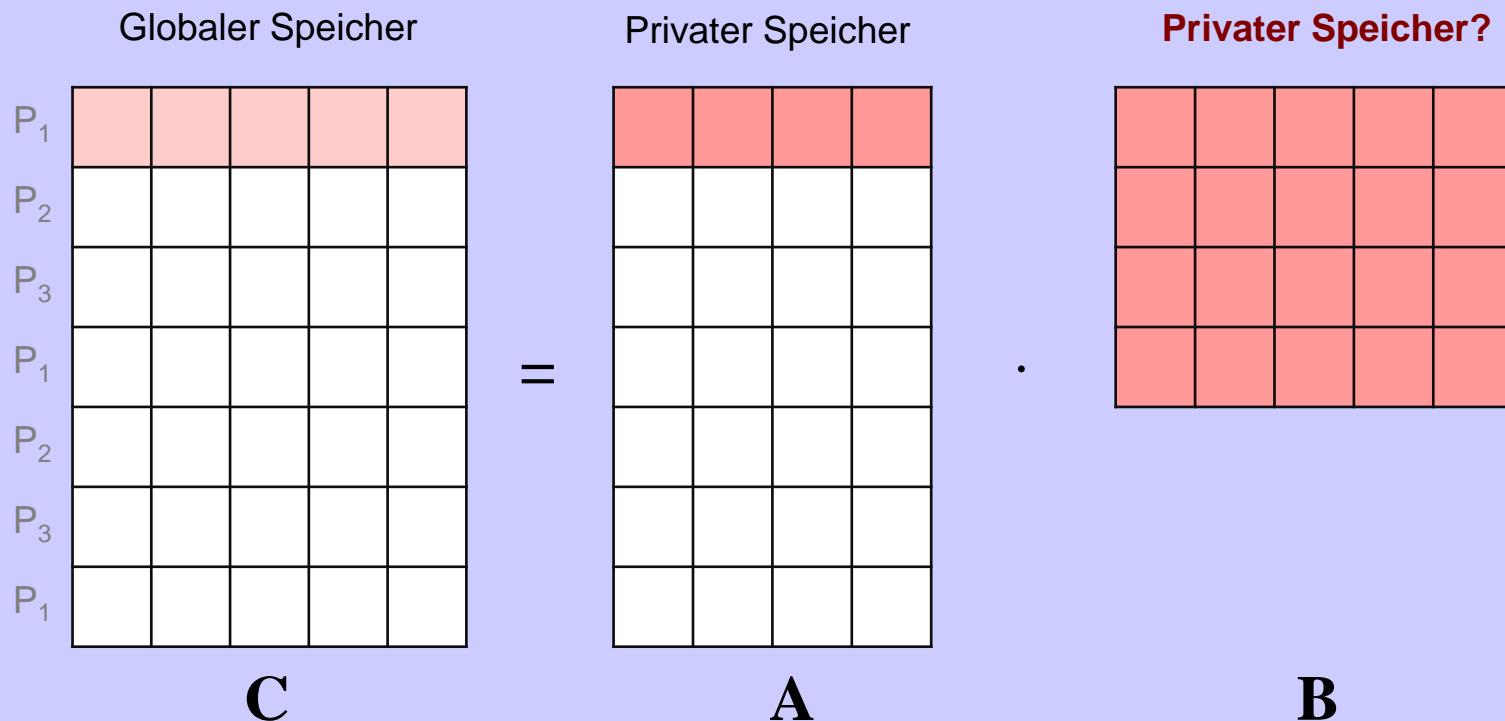


# Version 5: C zeilenweise, A privat

```
const char *KernelSource =  
"#define DIM 1000                                     // Groesse der Matrix \n"  
"__kernel void matmult(__global float *A, __global float *B, __global float *C) { \n"  
"    int k, j, i = get_global_id(0);                  \n"  
"    float Al[DIM], sum;                            // Privater Zeilenpuffer \n"  
"    for (k = 0; k < DIM; k++)                      \n"  
"        Al[k] = A[i*DIM+k];                         // Zeile i von A in den Puffer kopieren \n"  
"    for (j = 0; j < DIM; j++) {                      \n"  
"        sum = 0.0;                                  \n"  
"        for (k = 0; k < DIM; k++)                    \n"  
"            sum += Al[k] * B[k*DIM+j];              \n"  
"        C[i*DIM+j] = sum;                          \n"  
"    }                                              \n"  
" }                                              \n"  
"\n";
```

```
void main(int argc, char **argv)  
{  
    size_t global[1] = {DIM};  
    // ...  
    clEnqueueNDRangeKernel(command_queue, kernel, 1, NULL, global, NULL, 0, NULL, NULL);
```

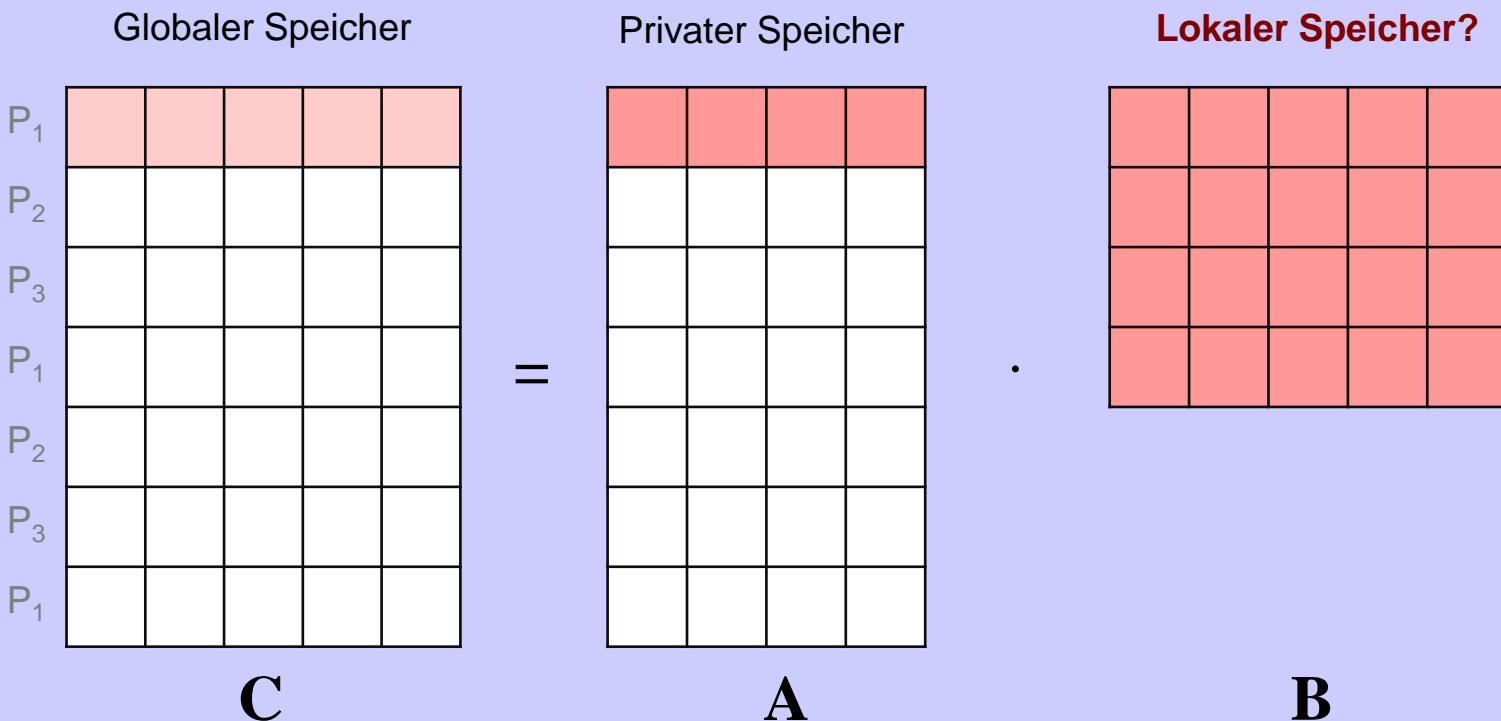
# Matrix-Multiplikation (zeilenweise)



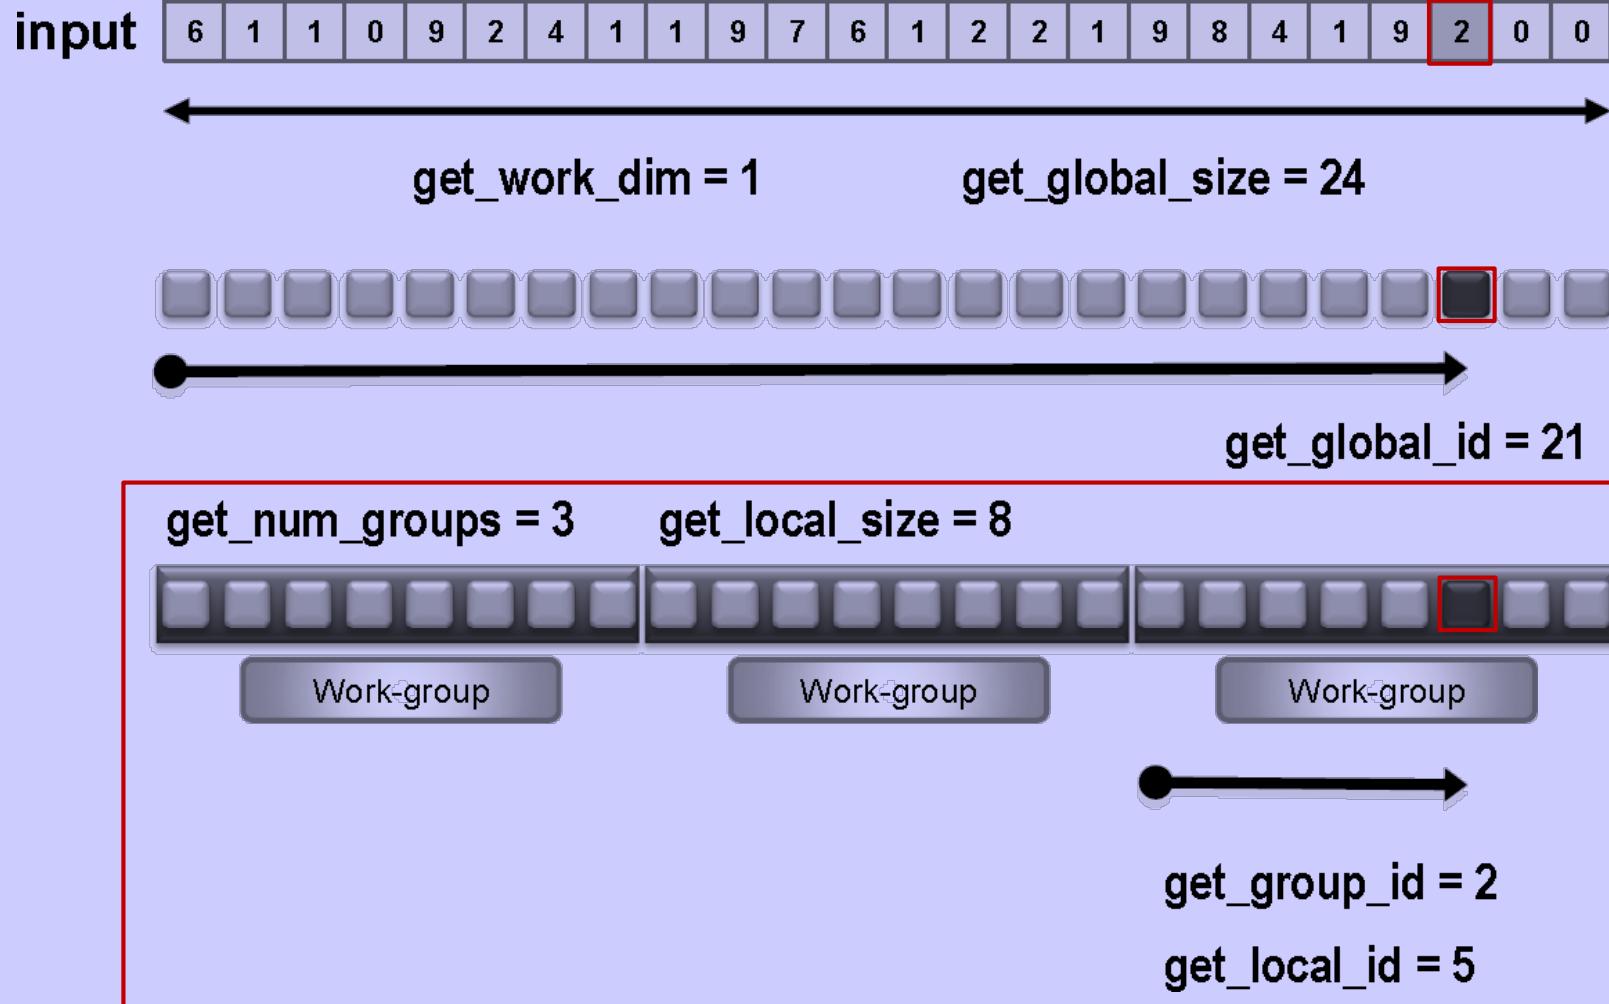
# Versuch 1: C zeilenweise, A und B privat

```
const char *KernelSource =  
"#define DIM 1000                                     // Groesse der Matrix \n"  
"__kernel void matmult(__global float *A, __global float *B, __global float *C) { \n"  
"    int k, j, i = get_global_id(0);  
"    float Al[DIM], Bl[DIM], sum;  
"    for (k = 0; k < DIM; k++)  
"        Al[k] = A[i*DIM+k];  
"    for (j = 0; j < DIM; j++) {  
"        for (k = 0; k < DIM; k++)  
"            Bl[k] = B[k*DIM+j];  
"        sum = 0.0;  
"        for (k = 0; k < DIM; k++)  
"            sum += Al[k] * Bl[k];  
"        C[i*DIM+j] = sum;  
"    }  
" }  
"\n";
```

# Matrix-Multiplikation (zeilenweise)



# Lokaler Speicher und Arbeitsgruppen



# Versuch 2: C zeilenweise, A priv. B lokal

```
const char *KernelSource =  
"#define DIM 1000                                     // Groesse der Matrix \n"  
"__kernel void matmult(__global float *A, __global float *B, __global float *C,    \n"  
"                      __local float *Bl) {                                \n"  
"    int k, j, i = get_global_id(0);                                \n"  
"    float Al[DIM], sum;                                         \n"  
"    int il = get_local_id(0);          // n work-items teilen sich das Kopieren \n"  
"    int nl = get_local_size(0);                                \n"  
"    for (k = 0; k < DIM; k++)                                \n"  
"        Al[k] = A[i*DIM+k];                                \n"  
"    for (j = 0; j < DIM; j++) {                                \n"  
"        for (k = il; k < DIM; k += nl)      // nur jede n-te Zeile von B kopieren \n"  
"            Bl[k] = B[k*DIM+j];                                \n"  
"        sum = 0.0;                                         \n"  
"        for (k = 0; k < DIM; k++)                                \n"  
"            sum += Al[k] * Bl[k];                                \n"  
"        C[i*DIM+j] = sum;                                \n"  
"    }                                              \n"  
"}
```

# Versuch 2: C zeilenweise, A priv. B lokal

```
const char *KernelSource =  
"#define DIM 1000                                     // Groesse der Matrix \n"  
"__kernel void matmult(__global float *A, __global float *B, __global float *C,    \n"  
"                      __local float *Bl) {           \n"  
"      // ...                                         \n"  
" }                                                 \n"  
\n";
```

**Explizite Definition der Work-group Größe (ganzzahlig teilbar):**

```
void main(int argc, char **argv)  
{  
    size_t global[1] = {DIM}, local[1];  
    // ...  
    clGetDeviceInfo(device_id, CL_DEVICE_MAX_COMPUTE_UNITS, sizeof(size_t), local, NULL);  
    local[0] = ggt(global[0], local[0]); // Groesster gemeinsamer Teiler ggt(1000, 24) = 8  
    // ...  
    clSetKernelArg(kernel, 3, sizeof(float)*DIM, NULL);  
    clEnqueueNDRangeKernel(command_queue, kernel, 1, NULL, global, local, 0, NULL, NULL);
```

# Größter gemeinsamer Teiler

```
int ggt(int x, int y)
{
    int z;

    while (y) {
        z = x % y;
        x = y;
        y = z;
    }
    return x;
}
```

## Beispiel:

$$\begin{array}{ccccccc} & x & \quad y & \quad & z \\ 1000 & / & 24 & = & 41 & \text{Rest} & 16 \\ & 24 & / & 16 & = & 1 & \text{Rest} & 8 \\ & 16 & / & 8 & = & 2 & \text{Rest} & 0 \\ & & & & & & \downarrow \end{array}$$

# Version 6: C zeilenweise, A priv. B lokal

```
const char *KernelSource =  
"#define DIM 1000                                     // Groesse der Matrix \n"  
"__kernel void matmult(__global float *A, __global float *B, __global float *C,    \n"  
"                      __local float *Bl) {                                \n"  
"    int k, j, i = get_global_id(0);                                \n"  
"    float Al[DIM], sum;                                         \n"  
"    int il = get_local_id(0);                                         \n"  
"    int nl = get_local_size(0);                                         \n"  
"    for (k = 0; k < DIM; k++)                                         \n"  
"        Al[k] = A[i*DIM+k];                                         \n"  
"    for (j = 0; j < DIM; j++) {                                         \n"  
"        for (k = il; k < DIM; k+=nl)                                         \n"  
"            Bl[k] = B[k*DIM+j];                                         \n"  
"        barrier(CLK_LOCAL_MEM_FENCE);          // Warten, dass alle B kopiert haben \n"  
"        sum = 0.0;                                         \n"  
"        for (k = 0; k < DIM; k++)                                         \n"  
"            sum += Al[k] * Bl[k];                                         \n"  
"        C[i*DIM+j] = sum;                                         \n"  
"    }                                         \n"  
"}
```

# Synchronisation und Arbeitsgruppen

`get_num_groups = 3`



`get_local_size = 8`



`get_group_id = 2`

`get_local_id = 5`

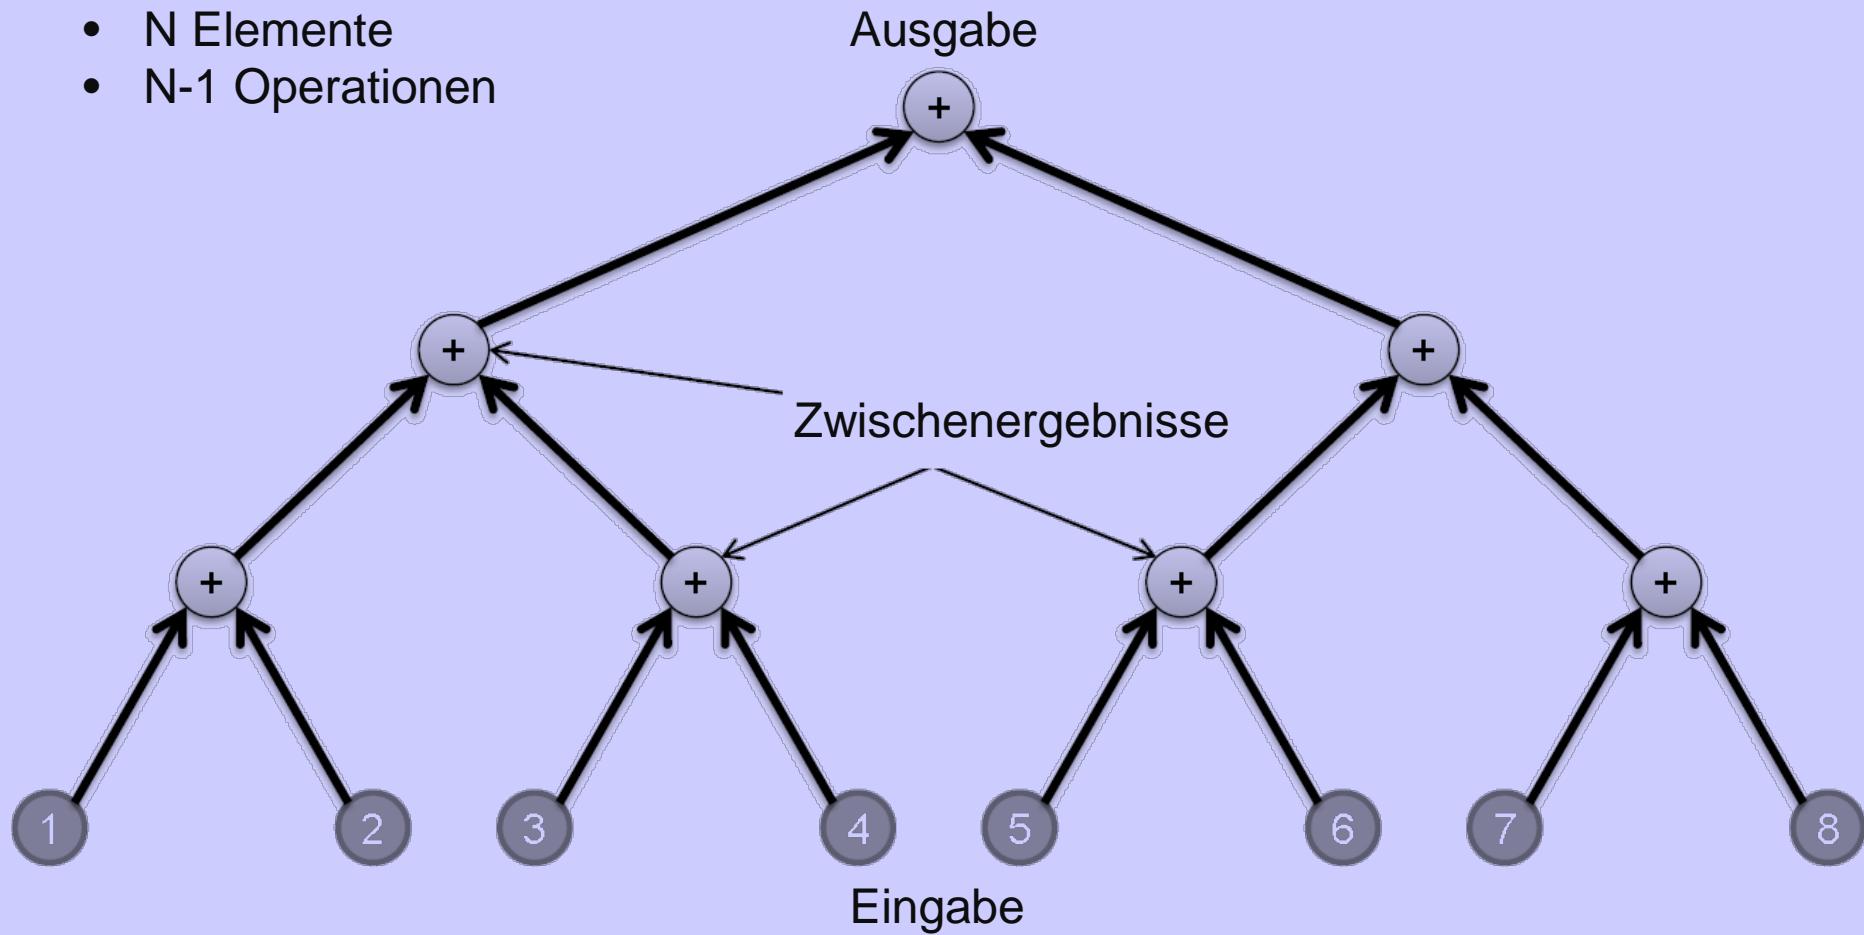
# Beispiel für Synchronisation

- **Ziel:** Reduziere eine Zahlenmenge auf einen einzelnen Wert
- **Beispiel:** Finde die Summe aller Elemente in einem Feld
- **Sequentieller Code**

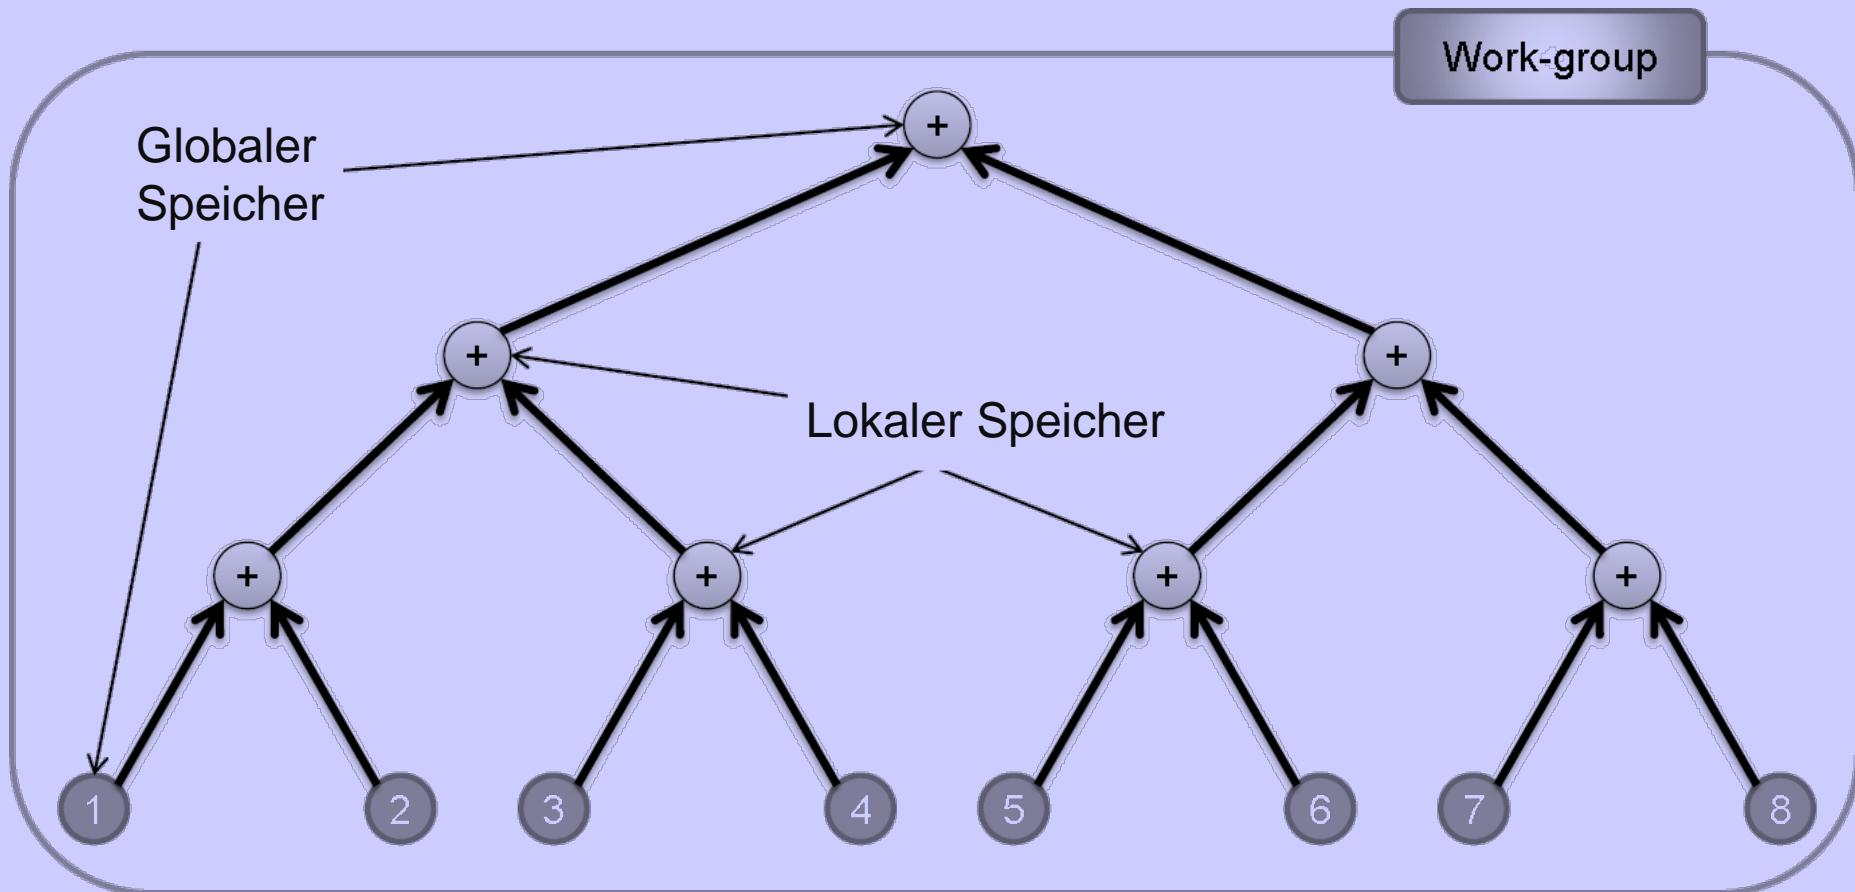
```
int reduce(int *A, int dim)
{
    int sum = 0;
    for (int i=0; i<dim; i++)
        sum += A[i];
    return sum;
}
```

# Parallele Reduktion

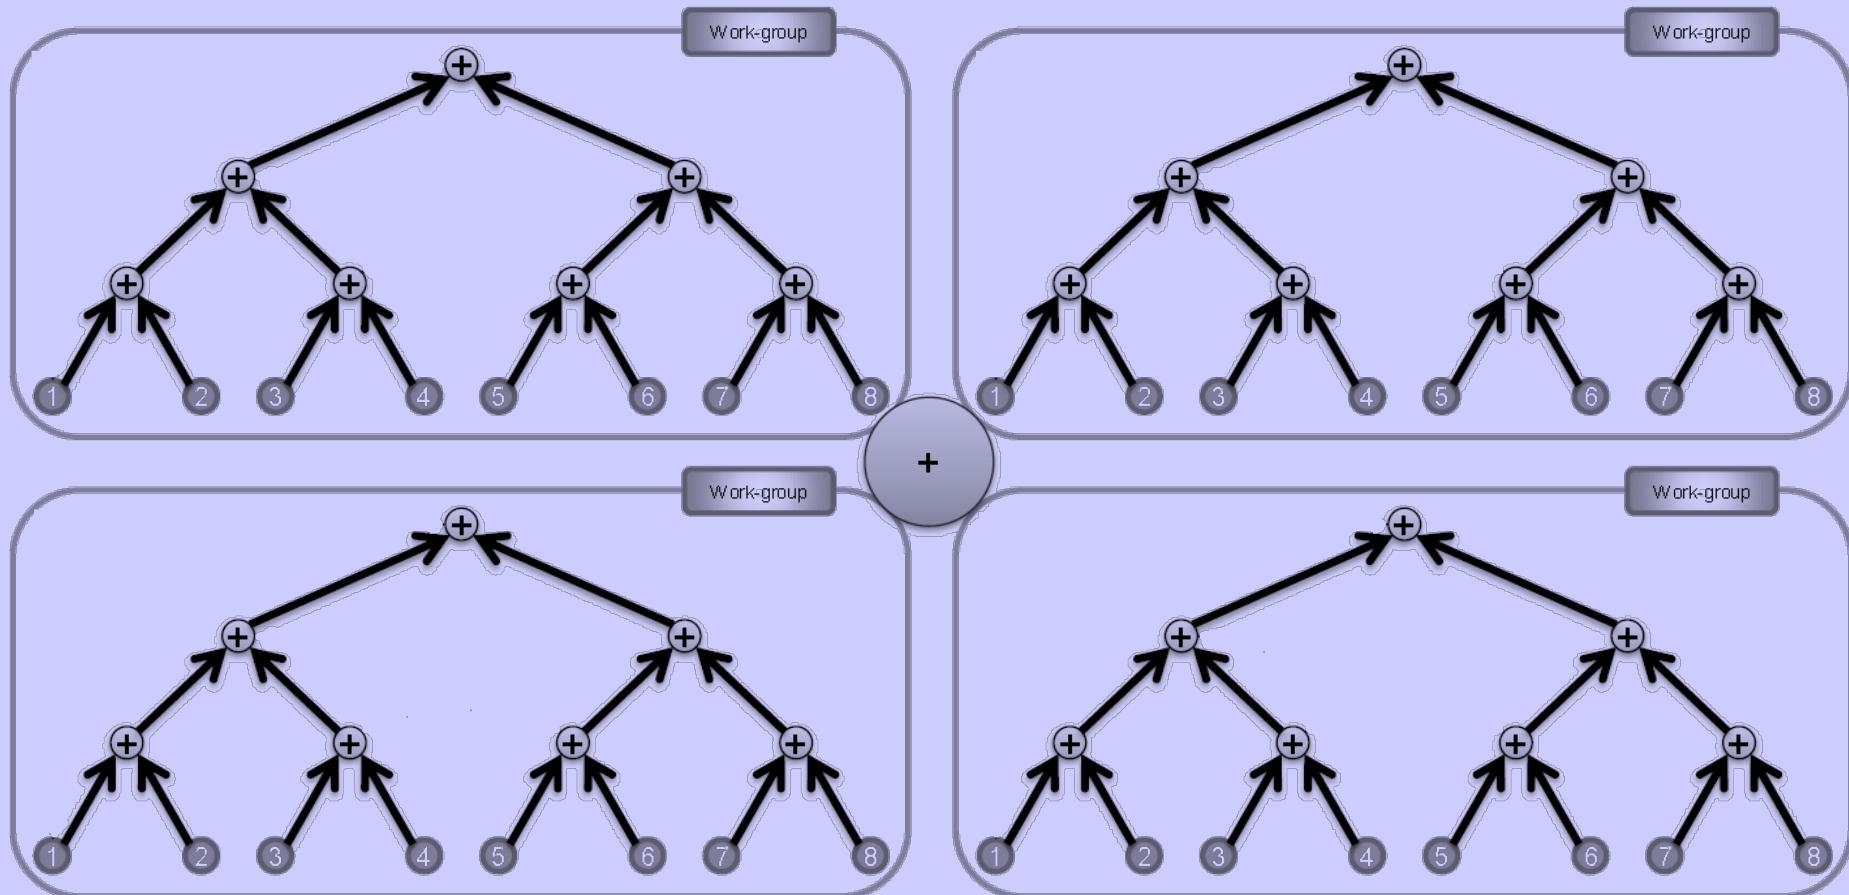
- N Elemente
- N-1 Operationen



# Parallele Reduktion in OpenCL



# Reduktion in OpenCL



# OpenCL Code 1

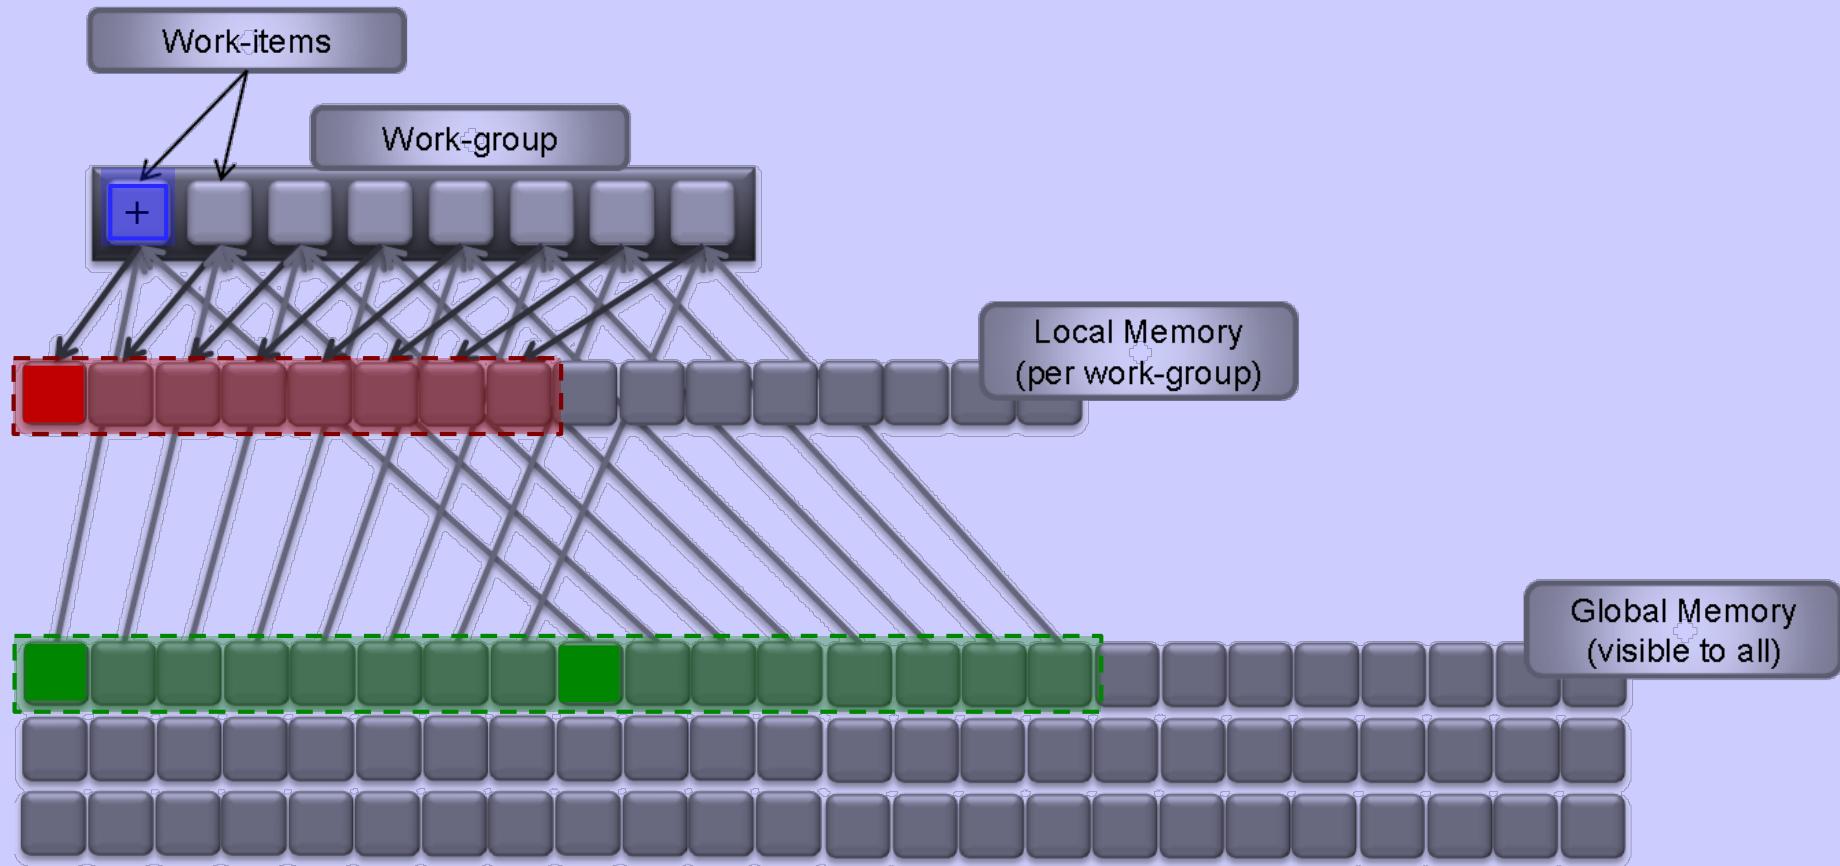
```
__kernel void
reduce(__global float *input,
         __global float *output,
         __local float *lmem)
{
    int gsize = get_global_size(0);
    int lsize = get_local_size(0);

    int gid    = get_global_id(0);
    int lid    = get_local_id(0);
    int oid    = get_group_id(0);

    // ...
}
```

```
// 1. Daten in lokalen Speicher kopieren
lmem[lid] = input[gid] + input[gid + gsize];
barrier(CLK_LOCAL_MEM_FENCE);
```

# Schritt 1: In lokalen Speicher kopieren



# OpenCL Code 2

```
__kernel void
reduce(__global float *input,
         __global float *output,
         __local float *lmem)
{
    int gsize = get_global_size(0);
    int lsize = get_local_size(0);

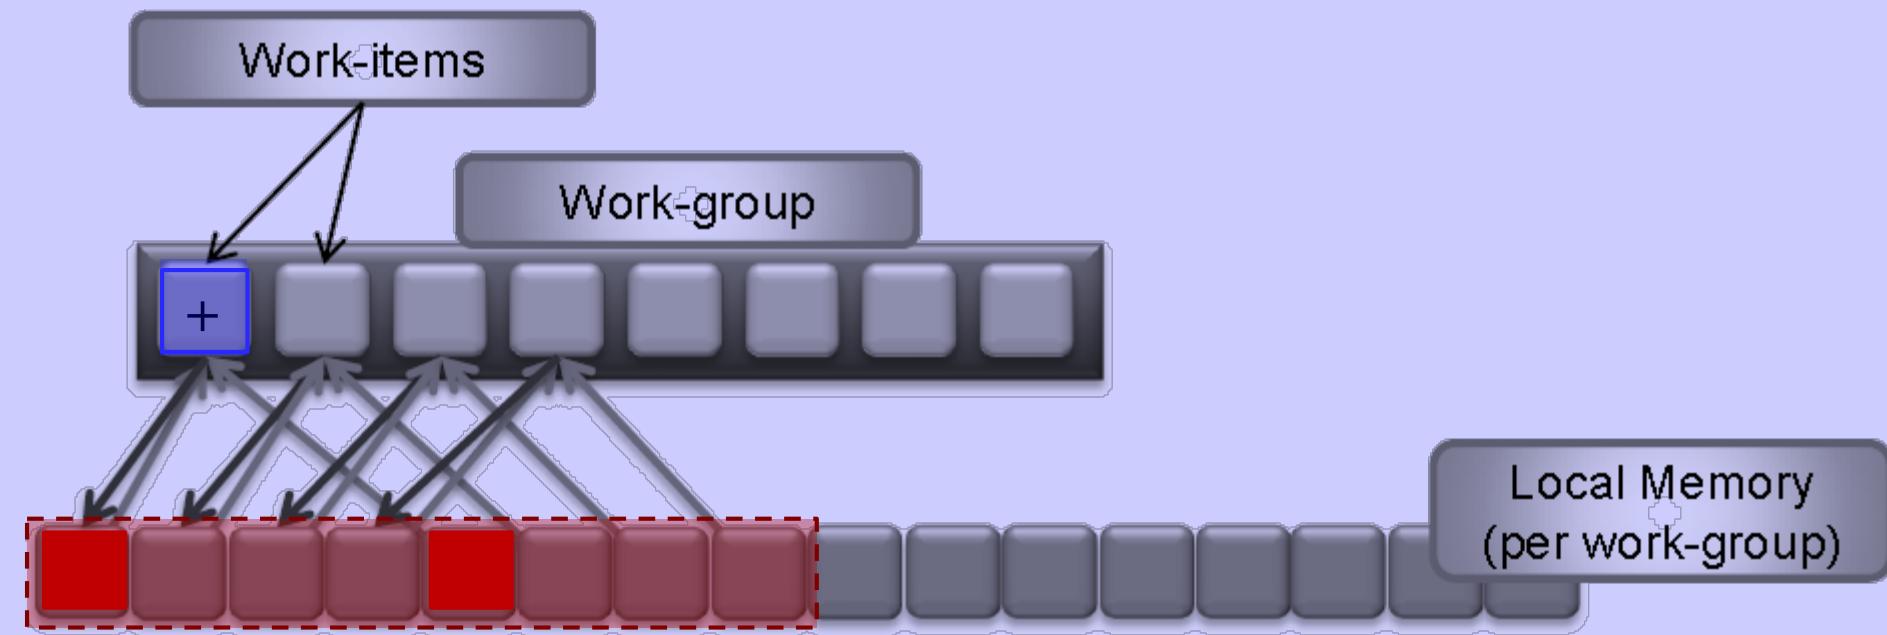
    int gid    = get_global_id(0);
    int lid    = get_local_id(0);
    int oid    = get_group_id(0);

    // ...
}
```

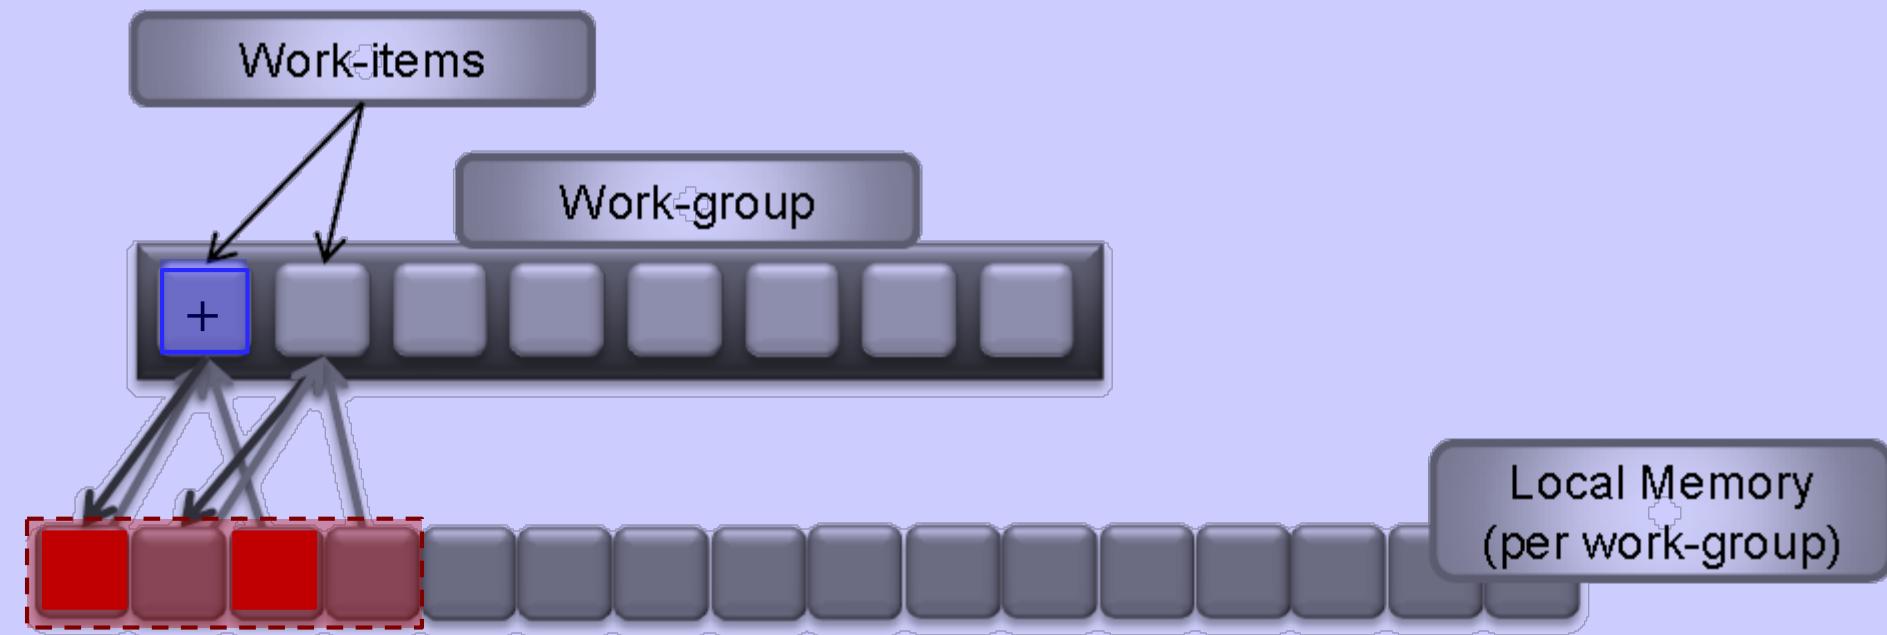
```
// 1. Daten in lokalen Speicher kopieren
lmem[lid] = input[gid] + input[gid + gsize];
barrier(CLK_LOCAL_MEM_FENCE);
```

```
// 2. Wiederhole lokale Reduktion
for (int s = lsize/2; s > 1; s /= 2) {
    if (lid < s)
        lmem[lid] += lmem[lid + s];
    barrier(CLK_LOCAL_MEM_FENCE);
}
```

# Schritt 2: Lokale Reduktion



# Schritt 2: Wiederholte Reduktion



# OpenCL Code 3

```
__kernel void
reduce(__global float *input,
       __global float *output,
       __local float *lmem)
{
    int gsize = get_global_size(0);
    int lsize = get_local_size(0);

    int gid    = get_global_id(0);
    int lid    = get_local_id(0);
    int oid    = get_group_id(0);

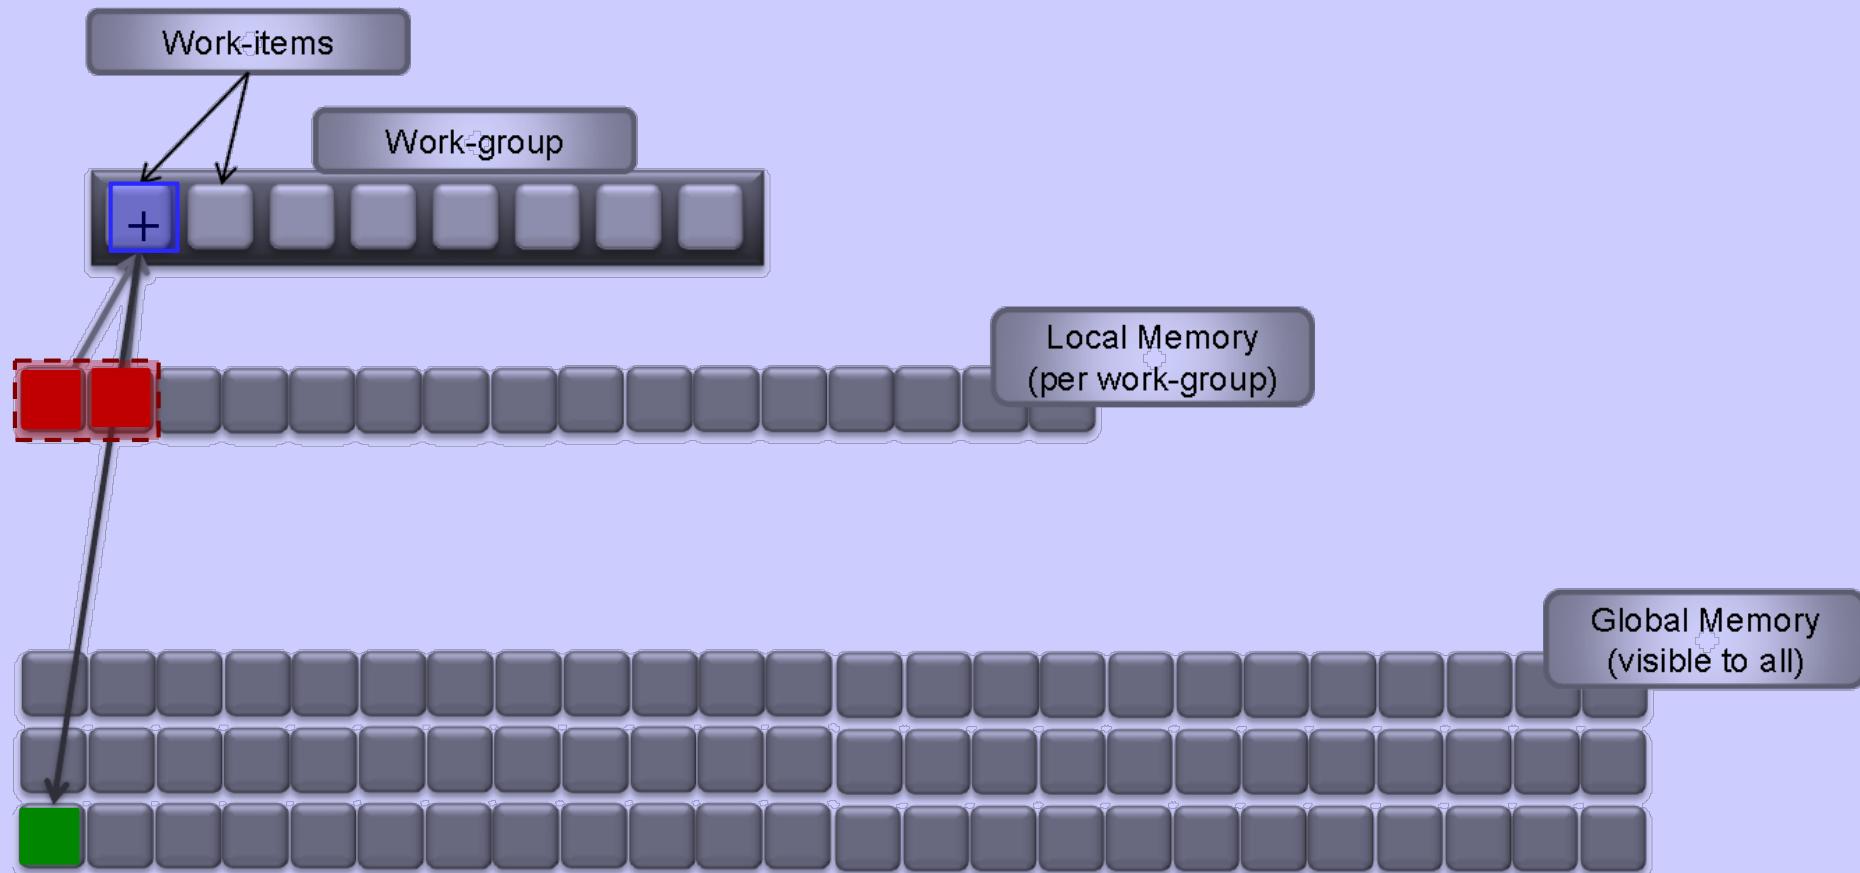
    // ...
}
```

```
// 1. Daten in lokalen Speicher kopieren
lmem[lid] = input[gid] + input[gid + gsize];
barrier(CLK_LOCAL_MEM_FENCE);
```

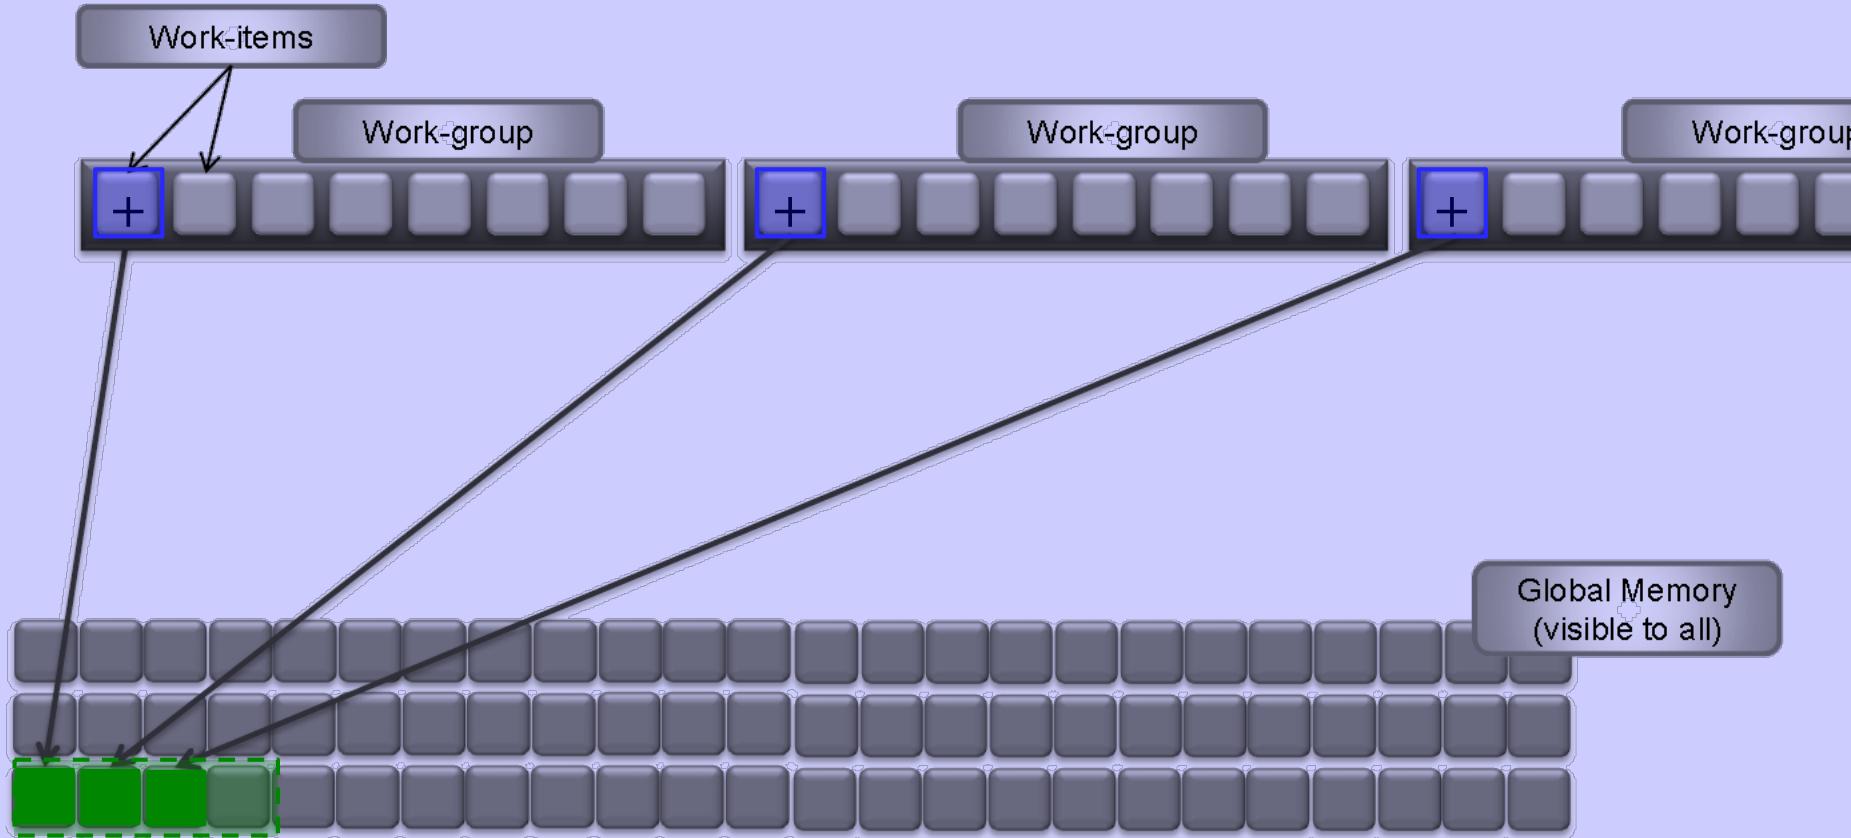
```
// 2. Wiederhole lokale Reduktion
for (int s = lsize/2; s > 1; s /= 2) {
    if (lid < s)
        lmem[lid] += lmem[lid + s];
    barrier(CLK_LOCAL_MEM_FENCE);
}
```

```
// 3. Schreibe Ergebnis in globalen Speicher
if (lid == 0)
    output[oid] = lmem[0] + lmem[1];
```

# Schritt 3: Ergebnis zurückkopieren



# Schritt 3: Alle Teilergebnisse



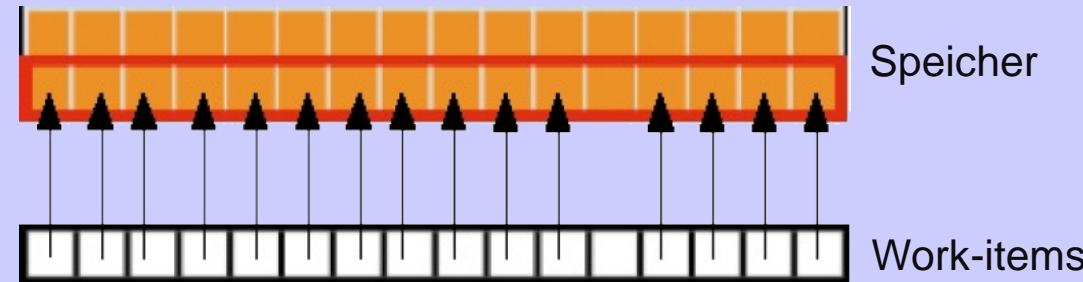
Die einzelnen Teilergebnisse der Work-groups mittels CPU aufsummieren oder den Kernel noch einmal innerhalb einer Work-group laufen lassen!



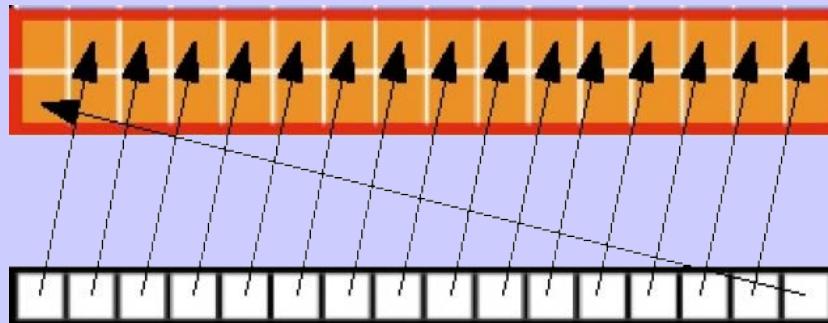
# Speicherorganisation

# Datenreihenfolge im Speicher

Alle Work-items können gleichzeitig auf Speicher zugreifen



Falsch ausgerichteter Speicher

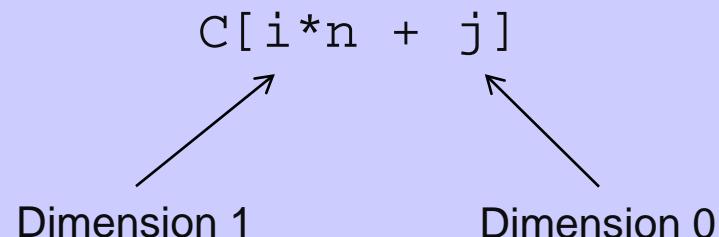
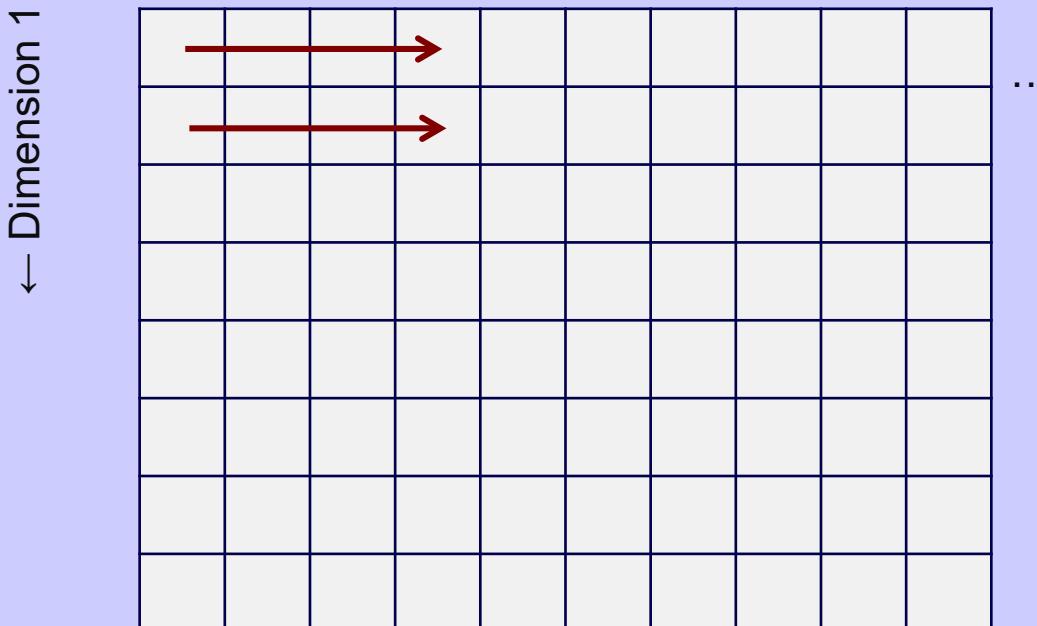


```
int i = get_global_id(0) + offset;  
out[i] = in[i];
```

```
int i = get_global_id(0) * stride;  
out[i] = in[i];
```

# Matrix Datenreihenfolge

Dimension 0 →



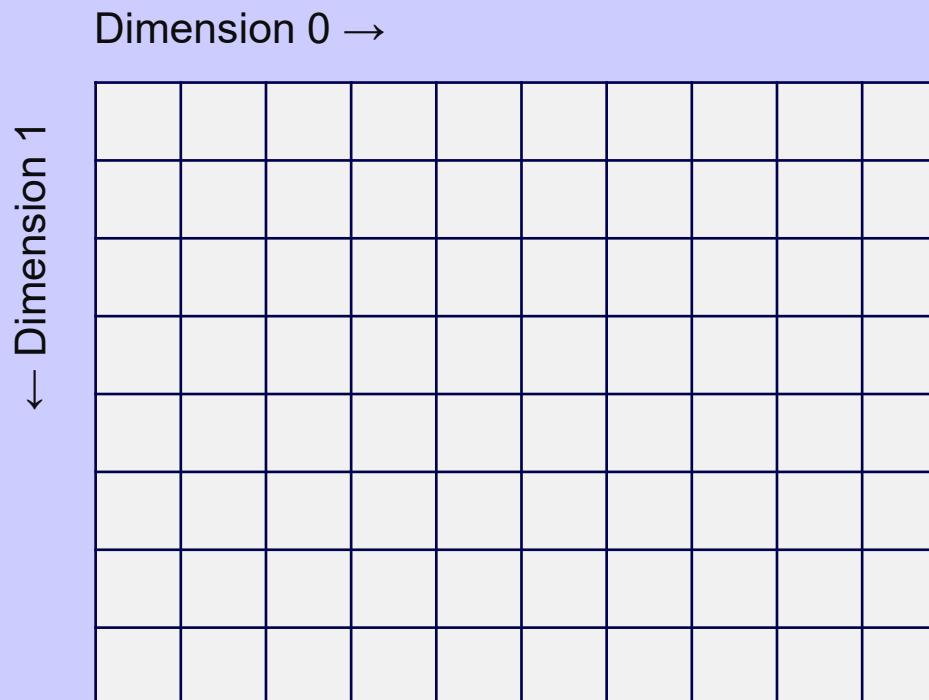
# NDRange Kernel Offset

- Bestimme Topologie (NDRange) und führe Kernel aus
- Parameter:

```
cl_int clEnqueueNDRangeKernel(  
    cl_command_queue command_queue,  
    cl_kernel kernel,  
    cl_uint work_dim,  
    const size_t *global_work_offset,  
    const size_t *global_work_size,  
    const size_t *local_work_size,  
    cl_uint num_events_in_wait_list,  
    const cl_event *event_wait_list,  
    cl_event event)
```

# Beispiel: Aufteilung des NDRange

- $\text{work\_dim} = 2$ ,  $\text{global\_work\_size} = \{ 10, 8 \}$



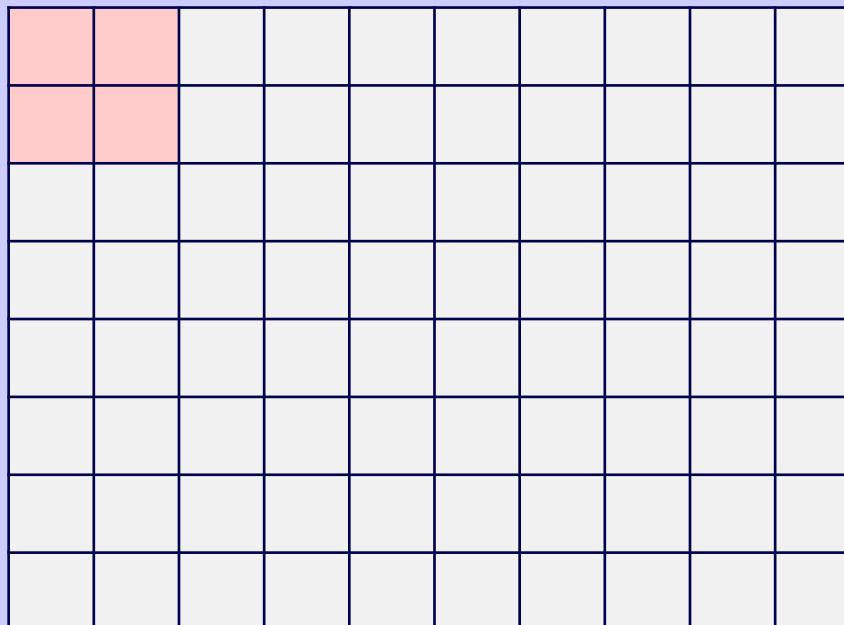
# NDRange Kernel Offset

- Bestimme Topologie (NDRange) und führe Kernel aus
- Parameter:

```
cl_int clEnqueueNDRangeKernel(  
    cl_command_queue command_queue,  
    cl_kernel kernel,  
    cl_uint work_dim,  
    const size_t global_work_offset,  
    const size_t global_work_size,  
    const size_t local_work_size,  
    cl_uint num_events_in_wait_list,  
    const cl_event *event_wait_list,  
    cl_event event)
```

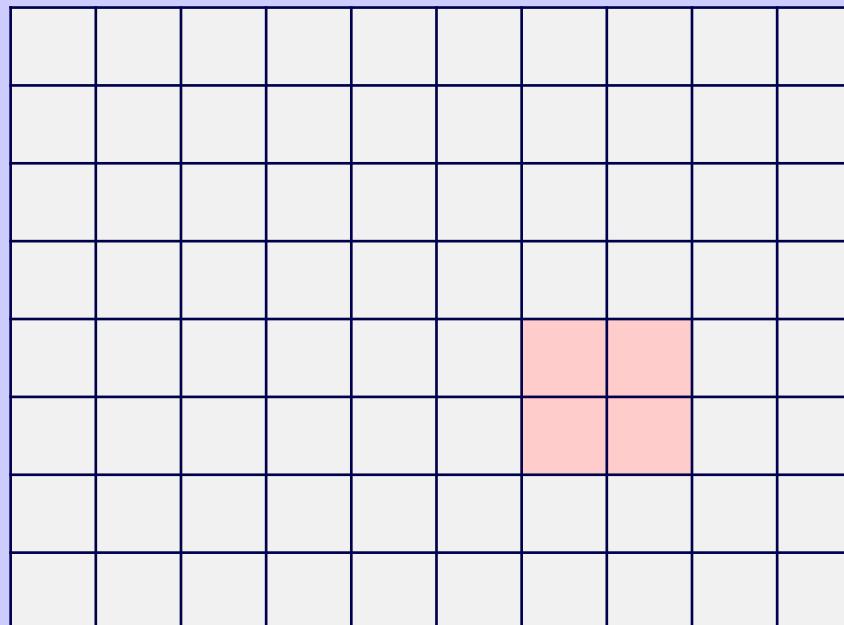
# Beispiel: Aufteilung des NDRange

- $\text{offset} = \{ 0, 0 \}$ ,  $\text{work\_size} = \{ 2, 2 \}$



# Beispiel: Aufteilung des NDRange

- $\text{offset} = \{ 6, 4 \}$ ,  $\text{work\_size} = \{ 2, 2 \}$



# Aufteilen der Kernel-Aufrufe

```
#define PART 10                                // DIM wird in PART Teile aufgeteilt

size_t part[2] = {DIM/PART, DIM/PART}, off[2];      // Teilbereiche von global[]
unsigned int i, j;

for (j = 0; j < PART; j++) {
    off[1] = j * part[1];
    for (i = 0; i < PART; i++) {
        off[0] = i * part[0];
        cEnqueueNDRangeKernel(command_queue, kernel, 2, off, part, NULL, 0, NULL, NULL);
    }
}
```

# Problem beim Aufteilen: Global Size?

```
__kernel void matmult(__global float *A, __global float *B, __global float *C, int n)
{
    int i, j, k;
    float sum = 0;

    i = get_global_id(0);
    j = get_global_id(1);

    for (k = 0; k < n; k++) {
        sum += A[i*n+k] * B[k*n+j];
    }
    C[i*n+j] = sum;
}
```

## Bemerkung:

- Dimension **n** der Matrix steht nicht mehr in **get\_global\_size()**, sondern muss als **#define** oder Funktionsargument übergeben werden.

# Problem beim Aufteilen: Zeitmessung?

```
#define PART 10                                     // DIM wird in PART Teile aufgeteilt

cl_event event1, event2;                           // Events für Zeitmessung
cl_int num = DIM;
size_t part[2] = {DIM/PART, DIM/PART}, off[2];    // Teilbereiche von global[]
unsigned int i, j;

clSetKernelArg(kernel, 3, sizeof(cl_int), &num);   // Matrixdimension als Parameter

for (j = 0; j < PART; j++) {
    off[1] = j * part[1];
    for (i = 0; i < PART; i++) {
        off[0] = i * part[0];
        clEnqueueNDRangeKernel(command_queue, kernel, 2, off, part, NULL, 0, NULL, &event2)
    }
    if ((i == 0) && (j == 0)) event1 = event2;           // Ersten Event merken
}
clFinish(command_queue);

clGetEventProfilingInfo(event1, CL_PROFILING_COMMAND_START, sizeof(start), &start, NULL);
clGetEventProfilingInfo(event2, CL_PROFILING_COMMAND_END, sizeof(end), &end, NULL);
printf("gpu: %.1f ms\n", ((end - start) * 1.0e-6));
```

# Weitere Konzepte in OpenCL

# Weitere Datentypen: Vektoren

char <i>n</i>	Ein Vektor von n 8-bit signed integer Werten
uchar <i>n</i>	Ein Vektor von n 8-bit unsigned integer Werten
short <i>n</i>	Ein Vektor von n 16-bit signed integer Werten
ushort <i>n</i>	Ein Vektor von n 16-bit unsigned integer Werten
int <i>n</i>	Ein Vektor von n 32-bit signed integer Werten
uint <i>n</i>	Ein Vektor von n 32-bit unsigned integer Werten
long <i>n</i>	Ein Vektor von n 64-bit signed integer Werten
ulong <i>n</i>	Ein Vektor von n 64-bit unsigned integer Werten
float <i>n</i>	Ein Vektor von n 32-bit floating-point Werten
double <i>n</i>	Ein Vektor von n 64-bit floating-point Werten
half <i>n</i>	Ein Vektor von n 16-bit floating-point Werten

*n* kann dabei 2, 3, 4, 8 oder 16 sein

# Doppelte Genauigkeit

Fließkommazahl mit doppelter Genauigkeit (**double**, 8Byte)

```
#if defined(cl_khr_fp64) // Khronos Erweiterung verfügbar?  
#pragma OPENCL EXTENSION cl_khr_fp64 : enable  
#elseif defined(cl_amd_fp64) // AMD Erweiterung verfügbar?  
#pragma OPENCL EXTENSION cl_amd_fp64 : enable  
#endif
```

# Vektordatentypen

```
// Ein Vektor vom Typ float4 mit 4 Gleitkommazahlen
float4 a = (float4)(0.0, 1.0, 2.0, 3.0);

// Ein Vektor vom Typ float2 mit 2 Gleitkommazahlen
float2 b = (float2)(1.0, 2.0);

// Äquivalent zu (0.0, 0.0)
float2 c = (float2)(0.0);

// Kombination von zwei float2 zu einem float4
float4 d = (float4)(b, c);
```

# Arithmetische Operationen

```
// Deklaration der Variablen
float4 a, b, c;
int8 d, e, f, g;

// Komponentenweise Addition von zwei float4 Vektoren
c = a + b;

// Komponentenweise Multiplikation von zwei int8 Vektoren
f = d * e;

// Multiplikation aller Komponenten eines int8 Vektors mit einer Zahl
g = 2 * f;
```

# Zugriff auf die Vektorkomponenten

	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
<b>float2 v;</b>	v.x, v.s0	v.y, v.s1														
<b>float4 v;</b>	v.x, v.s0	v.y, v.s1	v.z, v.s2	v.w, v.s3												
<b>float8 v;</b>	v.s0	v.s1	v.s2	v.s3	v.s4	v.s5	v.s6	v.s7								
<b>float16 v;</b>	v.s0	v.s1	v.s2	v.s3	v.s4	v.s5	v.s6	v.s7	v.s8	v.s9	v.sa, v.sA	v.sb, v.sB	v.sc, v.sC	v.sd, v.sD	v.se, v.sE	v.sf, v.sF

## Beispiele:

```
float2 pos;
pos.x = 1.0f;
pos.z = 1.0f; // ist illegal!

float4 c;
c.z = 1.0f;
c.xy = (float2)(3.0f, 4.0f);
c.xyzw = (float4)(1.0f, 2.0f, 3.0f, 4.0f);

float4 pos = (float4)(1.0f, 2.0f, 3.0f, 4.0f);
float4 swiz = pos.wzyx; // (4, 3, 2, 1)
float4 dup = pos.xxyy; // (1, 1, 2, 2)
```

Das **Mischen** ist nicht erlaubt:

```
float4 f;
float4 A = f.xs123;
float4 B = f.s012w;
```

# Auswahl von Vektorkomponenten

	v.lo	v.hi	v.odd	v.even
<b>float2</b>	v.x, v.s0	v.y, v.s1	v.y, v.s1	v.x, v.s0
<b>float4</b>	v.s01, v.xy	v.s23, v.zw	v.s13, v.yw	v.s02, v.xz
<b>float8</b>	v.s0123	v.s4567	v.s1357	v.s0246
<b>float16</b>	v.s01234567	v.s89abcdef	v.s13579bdf	v.s02468ace

# Mathematik Funktionen

## Math Built-in Functions [6.11.2]

*T* is type float or floatn (or optionally double, doublen, or halfn). intn, uintn, and ulongn must be scalar when *T* is scalar. The symbol **HN** indicates that Half and Native variants are available by prepending “half\_” or “native\_” to the function name, as in half\_cos() and native\_cos(). Optional extensions enable double, doublen, and halfn types.

<i>T</i> acos ( <i>T</i> )	Arc cosine
<i>T</i> acosh ( <i>T</i> )	Inverse hyperbolic cosine
<i>T</i> acospi ( <i>T</i> x)	acos ( <i>x</i> ) / $\pi$
<i>T</i> asin ( <i>T</i> )	Arc sine
<i>T</i> asinh ( <i>T</i> )	Inverse hyperbolic sine
<i>T</i> asinpi ( <i>T</i> x)	asin ( <i>x</i> ) / $\pi$
<i>T</i> atan ( <i>T</i> y_over_x)	Arc tangent
<i>T</i> atan2 ( <i>T</i> y, <i>T</i> x)	Arc tangent of <i>y</i> / <i>x</i>
<i>T</i> atanh ( <i>T</i> )	Hyperbolic arc tangent
<i>T</i> atanpi ( <i>T</i> x)	atan ( <i>x</i> ) / $\pi$
<i>T</i> atan2pi ( <i>T</i> x, <i>T</i> y)	atan2 ( <i>x</i> , <i>y</i> ) / $\pi$
<i>T</i> cbrt ( <i>T</i> )	cube root
<i>T</i> ceil ( <i>T</i> )	Round to integer toward + infinity
<i>T</i> copysign ( <i>T</i> x, <i>T</i> y)	<i>x</i> with sign changed to sign of <i>y</i>
<i>T</i> cos ( <i>T</i> ) <b>HN</b>	cosine
<i>T</i> cosh ( <i>T</i> )	hyperbolic consine
<i>T</i> cospi ( <i>T</i> x)	cos ( $\pi$ <i>x</i> )
<i>T</i> half_divide ( <i>T</i> x, <i>T</i> y)	<i>x</i> / <i>y</i>
<i>T</i> native_divide ( <i>T</i> x, <i>T</i> y)	
<i>T</i> erfc ( <i>T</i> )	Complementary error function
<i>T</i> erf ( <i>T</i> )	Calculates error function of <i>T</i>
<i>T</i> exp ( <i>T</i> x) <b>HN</b>	Exponential base e
<i>T</i> exp2 ( <i>T</i> ) <b>HN</b>	Exponential base 2

<i>T</i> exp10 ( <i>T</i> )	<b>HN</b>	Exponential base 10	
<i>T</i> expm1 ( <i>T</i> x)		$e^x - 1.0$	
<i>T</i> fabs ( <i>T</i> )		Absolute value	
<i>T</i> fdim ( <i>T</i> x, <i>T</i> y)		“Positive difference” between <i>x</i> and <i>y</i>	
<i>T</i> floor ( <i>T</i> )		Round to integer toward - infinity	
<i>T</i> fma ( <i>T</i> a, <i>T</i> b, <i>T</i> c)		Multiply and add, then round	
<i>T</i> fmax ( <i>T</i> x, <i>T</i> y)		Return <i>y</i> if <i>x</i> < <i>y</i> , otherwise it returns <i>x</i>	
<i>T</i> halfn fmax ( <i>halfn</i> <i>x</i> , <i>half</i> )			
<i>T</i> fmin ( <i>T</i> x, <i>T</i> y)		Return <i>y</i> if <i>y</i> < <i>x</i> , otherwise it returns <i>x</i>	
<i>T</i> halfn fmin ( <i>halfn</i> <i>x</i> , <i>half</i> )			
<i>T</i> fmod ( <i>T</i> x, <i>T</i> y)		Modulus. Returns <i>x</i> - <i>y</i> * trunc ( <i>x/y</i> )	
<i>T</i> fract ( <i>T</i> x, <i>T</i> *iptr)		Fractional value in <i>x</i>	
<i>T</i> frexp ( <i>T</i> x, intn *exp)		Extract mantissa and exponent	
<i>T</i> hypot ( <i>T</i> x, <i>T</i> y)		square root of $x^2 + y^2$	
<i>T</i> intn ilogb ( <i>T</i> x)		Return exponent as an integer value	
<i>T</i> ldexp ( <i>T</i> x, intn n)		$x * 2^n$	
<i>T</i> ldexp ( <i>T</i> x, int n)			
<i>T</i> lgamma ( <i>T</i> x)		Log gamma function	
<i>T</i> lgamma_r ( <i>T</i> x, intn *signp)			
<i>T</i> log ( <i>T</i> ) <b>HN</b>		Natural logarithm	
<i>T</i> log2 ( <i>T</i> ) <b>HN</b>		Base 2 logarithm	
<i>T</i> log10 ( <i>T</i> ) <b>HN</b>		Base 10 logarithm	
<i>T</i> log1p ( <i>T</i> x)		$\ln(1.0 + x)$	
<i>T</i> logb ( <i>T</i> x)		exponent of <i>x</i>	
<i>T</i> mad ( <i>T</i> a, <i>T</i> b, <i>T</i> c)		Approximates <i>a</i> * <i>b</i> + <i>c</i>	
<i>T</i> modf ( <i>T</i> x, <i>T</i> *iptr)		Decompose a floating-point number	
<i>T</i> nan ( <i>T</i> )		float nan (uintn nancode) floatn nan (uintn nancode) halfn nan (ushortn nancode) doublen nan (ulongn nancode) doublen nan (uintn nancode)	Quiet NaN
<i>T</i> nextafter ( <i>T</i> x, <i>T</i> y)			Next representable floating-point value following <i>x</i> in the direction of <i>y</i>
<i>T</i> pow ( <i>T</i> x, <i>T</i> y)			Compute <i>x</i> to the power of <i>y</i> ( $x^y$ )
<i>T</i> pown ( <i>T</i> x, intn y)			Compute $x^y$ , where <i>y</i> is an integer
<i>T</i> powr ( <i>T</i> x, <i>T</i> y) <b>HN</b>			Compute $x^y$ , where <i>x</i> is $\geq 0$
<i>T</i> half_recip ( <i>T</i> x)			$1/x$
<i>T</i> native_recip ( <i>T</i> x)			
<i>T</i> remainder ( <i>T</i> x, <i>T</i> y)			Floating point remainder function
<i>T</i> remquo ( <i>T</i> x, <i>T</i> y, intn *quo)			Floating point remainder and quotient function
<i>T</i> rint ( <i>T</i> )			Round integer to nearest even integer
<i>T</i> rootn ( <i>T</i> x, intn y)			Compute <i>x</i> to the power of $1/y$
<i>T</i> round ( <i>T</i> x)			Integral value nearest to <i>x</i> rounding
<i>T</i> rsqrt ( <i>T</i> ) <b>HN</b>			Inverse square root
<i>T</i> sin ( <i>T</i> ) <b>HN</b>			sine
<i>T</i> sincos ( <i>T</i> x, <i>T</i> *cosval)			sine and cosine of <i>x</i>
<i>T</i> sinh ( <i>T</i> )			hyperbolic sine
<i>T</i> sinpi ( <i>T</i> x)			$\sin(\pi x)$
<i>T</i> sqrt ( <i>T</i> ) <b>HN</b>			square root
<i>T</i> tan ( <i>T</i> ) <b>HN</b>			tangent
<i>T</i> tanh ( <i>T</i> )			hyperbolic tangent
<i>T</i> tanpi ( <i>T</i> x)			$\tan(\pi x)$
<i>T</i> tgamma ( <i>T</i> )			gamma function
<i>T</i> trunc ( <i>T</i> )			Round to integer toward zero

# Beispiel: Neue Native-/Fast-Funktionen

```
kernel void regularFuncs()
{
    for (int i=0; i<5000; i++)
    {
        float a=1, b=2, c=3, d=4;
        float e = a*b+c;
        e = a*b+c*d;
        e = sin(a);
        e = cos(b);

        float4 vec1 = (float4)(1, 2, 3, 0);
        float4 vec2 = (float4)(-1, 3, 1, 0);
        float4 vec = distance(vec1, vec2);

        double x=1, y=2, z=3;
        double resp = x*y+z;
    }
}
```

```
kernel void nativeFuncs()
{
    for (int i=0; i<5000; i++)
    {
        float a=1, b=2, c=3, d=4;
        float e = mad(a,b,c);
        e = mad(a,b,c*d);
        e = native_sin(a);
        e = native_cos(b);

        float4 vec1 = (float4)(1, 2, 3, 0);
        float4 vec2 = (float4)(-1, 3, 1, 0);
        float4 vec = fast_distance(vec1, vec2);

        double x=1, y=2, z=3;
        double resp = mad(x,y,z);
    }
}
```

# Geometrie Funktionen

## Geometric Built-in Functions [6.11.5]

Vector types may have 2 or 4 components. Optional extensions enable double, doublen, and halfn types.

float4 cross (float4 <i>p0</i> , float4 <i>p1</i> ) double4 cross (double4 <i>p0</i> , double4 <i>p1</i> ) half4 cross (half4 <i>p0</i> , half4 <i>p1</i> )	Cross product
float dot (float <i>p0</i> , float <i>p1</i> ) float dot (floatn <i>p0</i> , floatn <i>p1</i> ) double dot (double <i>p0</i> , double <i>p1</i> ) double dot (doublen <i>p0</i> , doublen <i>p1</i> ) half dot (half <i>p0</i> , half <i>p1</i> ) half dot (halfn <i>p0</i> , halfn <i>p1</i> )	Dot product
float distance (float <i>p0</i> , float <i>p1</i> ) float distance (floatn <i>p0</i> , floatn <i>p1</i> ) double distance (double <i>p0</i> , double <i>p1</i> ) double distance (doublen <i>p0</i> , doublen <i>p1</i> ) half distance (half <i>p0</i> , half <i>p1</i> ) half distance (halfn <i>p0</i> , halfn <i>p1</i> )	Vector distance
float length (float <i>p</i> ) float length (floatn <i>p</i> ) double length (double <i>p</i> ) double length (doublen <i>p</i> ) half length (half <i>p</i> ) half length (halfn <i>p</i> )	Vector length
float normalize (float <i>p</i> ) floatn normalize (floatn <i>p</i> ) double normalize (double <i>p</i> ) doublen normalize (doublen <i>p</i> ) half normalize (half <i>p</i> ) halfn normalize (halfn <i>p</i> )	Normal vector length 1
float fast_distance (float <i>p0</i> , float <i>p1</i> ) float fast_distance (floatn <i>p0</i> , floatn <i>p1</i> )	Vector distance
float fast_length (float <i>p</i> ) float fast_length (floatn <i>p</i> )	Vector length
float fast_normalize (float <i>p</i> ) floatn fast_normalize (floatn <i>p</i> )	Normal vector length 1

## Common Built-in Functions [6.11.4]

*T* is type float or floatn (or optionally double, doublen, or halfn). Optional extensions enable double, doublen, and halfn types.

<i>T clamp</i> ( <i>T x, T min, T max</i> ) floatn clamp (floatn <i>x</i> , float <i>min</i> , float <i>max</i> ) doublen clamp (doublen <i>x</i> , double <i>min</i> , double <i>max</i> ) halfn clamp (halfn <i>x</i> , half <i>min</i> , half <i>max</i> )	Clamp <i>x</i> to range given by <i>min, max</i>
<i>T degrees</i> ( <i>T radians</i> )	radians to degrees
<i>T max</i> ( <i>T x, T y</i> ) floatn max (floatn <i>x</i> , float <i>y</i> ) doublen max (doublen <i>x</i> , double <i>y</i> ) halfn max (halfn <i>x</i> , half <i>y</i> )	Max of <i>x</i> and <i>y</i>
<i>T min</i> ( <i>T x, T y</i> ) floatn min (floatn <i>x</i> , float <i>y</i> ) doublen min (doublen <i>x</i> , double <i>y</i> ) halfn min (halfn <i>x</i> , half <i>y</i> )	Min of <i>x</i> and <i>y</i>
<i>T mix</i> ( <i>T x, T y, T a</i> ) floatn mix (floatn <i>x</i> , float <i>y</i> , float <i>a</i> ) doublen mix (doublen <i>x</i> , double <i>y</i> , <i>a</i> ) halfn mix (halfn <i>x</i> , half <i>y</i> , <i>a</i> )	Linear blend of <i>x</i> and <i>y</i>
<i>T radians</i> ( <i>T degrees</i> )	degrees to radians
<i>T step</i> ( <i>T edge, Tx</i> ) floatn step (float <i>edge</i> , floatn <i>x</i> ) doublen step (double <i>edge</i> , doublen <i>x</i> ) halfn step (half <i>edge</i> , halfn <i>x</i> )	0.0 if <i>x</i> < <i>edge</i> , else 1.0
<i>T smoothstep</i> ( <i>T edge0, T edge1, Tx</i> ) floatn smoothstep (float <i>edge0</i> , float <i>edge1</i> , floatn <i>x</i> ) doublen smoothstep (double <i>edge0</i> , double <i>edge1</i> , doublen <i>x</i> ) halfn smoothstep (half <i>edge0</i> , half <i>edge1</i> , halfn <i>x</i> )	Step and interpolate
<i>T sign</i> ( <i>T x</i> )	Sign of <i>x</i>

# Vergleichsoperatoren 1

## Relational Built-in Functions [6.11.6]

*T* is type float, floatn, char, charn, uchar, ucharn, short, shortn, ushort, ushortn, int, intn, uint, uintn, long, longn, ulong, or ulongn (and optionally double, doublen). *S* is type char, charn, short, shortn, int, intn, long, or longn. *U* is type uchar, ucharn, ushort, ushortn, uint, uintn, ulong, or ulongn. Optional extensions enable double, doublen, and halfn types.

int <b>lsequal</b> (float <i>x</i> , float <i>y</i> ) intn <b>lsequal</b> (floatn <i>x</i> , floatn <i>y</i> ) int <b>lsequal</b> (double <i>x</i> , double <i>y</i> ) longn <b>lsequal</b> (doublen <i>x</i> , doublen <i>y</i> ) int <b>lsequal</b> (half <i>x</i> , half <i>y</i> ) shortn <b>lsequal</b> (halfn <i>x</i> , halfn <i>y</i> )	Compare of <i>x</i> == <i>y</i>
--	---------------------------------

int <b>lsnotequal</b> (float <i>x</i> , float <i>y</i> ) intn <b>lsnotequal</b> (floatn <i>x</i> , floatn <i>y</i> ) int <b>lsnotequal</b> (double <i>x</i> , double <i>y</i> ) longn <b>lsnotequal</b> (doublen <i>x</i> , doublen <i>y</i> ) int <b>lsnotequal</b> (half <i>x</i> , half <i>y</i> ) shortn <b>lsnotequal</b> (halfn <i>x</i> , halfn <i>y</i> )	Compare of <i>x</i> != <i>y</i>
--	---------------------------------

int <b>lsgreater</b> (float <i>x</i> , float <i>y</i> ) intn <b>lsgreater</b> (floatn <i>x</i> , floatn <i>y</i> ) int <b>lsgreater</b> (double <i>x</i> , double <i>y</i> ) longn <b>lsgreater</b> (doublen <i>x</i> , doublen <i>y</i> ) int <b>lsgreater</b> (half <i>x</i> , half <i>y</i> ) shortn <b>lsgreater</b> (halfn <i>x</i> , halfn <i>y</i> )	Compare of <i>x</i> > <i>y</i>
--	--------------------------------

int <b>lsgreaterequal</b> (float <i>x</i> , float <i>y</i> ) intn <b>lsgreaterequal</b> (floatn <i>x</i> , floatn <i>y</i> ) int <b>lsgreaterequal</b> (double <i>x</i> , double <i>y</i> ) longn <b>lsgreaterequal</b> (doublen <i>x</i> , doublen <i>y</i> ) int <b>lsgreaterequal</b> (half <i>x</i> , half <i>y</i> ) shortn <b>lsgreaterequal</b> (halfn <i>x</i> , halfn <i>y</i> )	Compare of <i>x</i> >= <i>y</i>
--	---------------------------------

int <b>Isinf</b> (float) intn <b>Isinf</b> (floatn) int <b>Isinf</b> (double) longn <b>Isinf</b> (doublen) int <b>Isinf</b> (half) shortn <b>Isinf</b> (halfn)	Test for +ve or -ve infinity
int <b>Isnan</b> (float) intn <b>Isnain</b> (floatn) int <b>Isnain</b> (double) longn <b>Isnain</b> (doublen) int <b>Isnain</b> (half) shortn <b>Isnain</b> (halfn)	Test for a NaN
int <b>Isnormal</b> (float) intn <b>Isnormal</b> (floatn) int <b>Isnormal</b> (double) longn <b>Isnormal</b> (doublen) int <b>Isnormal</b> (half) shortn <b>Isnormal</b> (halfn)	Test for a normal value
int <b>Isordered</b> (float <i>x</i> , float <i>y</i> ) intn <b>Isordered</b> (floatn <i>x</i> , floatn <i>y</i> ) int <b>Isordered</b> (double <i>x</i> , double <i>y</i> ) longn <b>Isordered</b> (doublen <i>x</i> , doublen <i>y</i> ) int <b>Isordered</b> (half <i>x</i> , half <i>y</i> ) shortn <b>Isordered</b> (halfn <i>x</i> , halfn <i>y</i> )	Test if arguments are ordered
int <b>Isunordered</b> (float <i>x</i> , float <i>y</i> ) intn <b>Isunordered</b> (floatn <i>x</i> , floatn <i>y</i> ) int <b>Isunordered</b> (double <i>x</i> , double <i>y</i> ) longn <b>Isunordered</b> (doublen <i>x</i> , doublen <i>y</i> ) int <b>Isunordered</b> (half <i>x</i> , half <i>y</i> ) shortn <b>Isunordered</b> (halfn <i>x</i> , halfn <i>y</i> )	Test if arguments are unordered

# Vergleichsoperatoren 2

<pre>int Isless (float x, float y) intn Isless (floatn x, floatn y) int Isless (double x, double y) longn Isless (doublen x, doublen y) int Isless (half x, half y) shortn Isless (halfn x, halfn y)</pre>	Compare of $x < y$	<pre>int signbit (float) intn signbit (floatn) int signbit (double) longn signbit (doublen) int signbit (half) shortn signbit (halfn)</pre>	Test for sign bit
<pre>int Islessequal (float x, float y) intn Islessequal (floatn x, floatn y) int Islessequal (double x, double y) longn Islessequal (doublen x, doublen y) int Islessequal (half x, half y) shortn Islessequal (halfn x, halfn y)</pre>	Compare of $x \leq y$	<pre>int any (S x)</pre>	1 if MSB in any component of $x$ is set; else 0
		<pre>int all (S x)</pre>	1 if MSB in all components of $x$ are set; else 0
<pre>int Islessgreater (float x, float y) intn Islessgreater (floatn x, floatn y) int Islessgreater (double x, double y) longn Islessgreater (doublen x, doublen y) int Islessgreater (half x, half y) shortn Islessgreater (halfn x, halfn y)</pre>	Compare of $(x < y) \mid\mid (x > y)$	<pre>T bitselect (Ta, Tb, Tc) halfn bitselect (halfna, halfnb, halfnc) doublen bitselect (doublena, doublenb, doublenc)</pre>	Each bit of result is corresponding bit of $a$ if corresponding bit of $c$ is 0
<pre>int Isfinite (float) intn Isfinite (floatn) int Isfinite (double) longn Isfinite (doublen) int Isfinite (half) shortn Isfinite (halfn)</pre>	Test for finite value	<pre>T select (Ta, Tb, Sc) T select (Ta, Tb, Uc) doublen select (doublen, doublen, longn)</pre>	For each component of a vector type, $\text{result}[i] = \text{if MSB of } c[i] \text{ is set ? } b[i] : a[i]$ For scalar type, $\text{result} = c ? b : a$

# 2D und 3D Bildobjekte

## Image Objects

### Create Image Objects [5.2.4]

```
cl_mem clCreateImage2D (
    cl_context context, cl_mem_flags flags,
    const cl_image_format *image_format,
    size_t image_width, size_t image_height,
    size_t image_row_pitch, void *host_ptr,
    cl_int *errcode_ret)

flags: CL_MEM_READ_WRITE,          CL_MEM_WRITE_ONLY,
CL_MEM_READ_ONLY,                 CL_MEM_USE_HOST_PTR,
CL_MEM_ALLOC_HOST_PTR,            CL_MEM_COPY_HOST_PTR
```

```
cl_mem clCreateImage3D (
    cl_context context, cl_mem_flags flags,
    const cl_image_format *image_format,
    size_t image_width, size_t image_height,
    size_t image_depth, size_t image_row_pitch,
    size_t image_slice_pitch, void *host_ptr,
    cl_int *errcode_ret)

flags: CL_MEM_READ_WRITE,          CL_MEM_WRITE_ONLY,
CL_MEM_READ_ONLY,                 CL_MEM_USE_HOST_PTR,
CL_MEM_ALLOC_HOST_PTR,            CL_MEM_COPY_HOST_PTR
```

### Query List of Supported Image Formats [5.2.5]

```
cl_int clGetSupportedImageFormats (
    cl_context context, cl_mem_flags flags,
    cl_mem_object_type image_type,
    cl_uint num_entries,
    cl_image_format *image_formats,
    cl_uint *num_image_formats)

flags: CL_MEM_READ_WRITE,          CL_MEM_WRITE_ONLY,
CL_MEM_READ_ONLY,                 CL_MEM_USE_HOST_PTR,
CL_MEM_ALLOC_HOST_PTR,            CL_MEM_COPY_HOST_PTR
```

### Copy Between Image and Buffer Objects [5.2.7]

```
cl_int clEnqueueCopyImageToBuffer (
    cl_command_queue command_queue,
    cl_mem src_image, cl_mem dst_buffer,
    const size_t src_origin[3], const size_t region[3],
    size_t dst_offset, cl_uint num_events_in_wait_list,
    const cl_event *event_wait_list, cl_event *event)

cl_int clEnqueueCopyBufferToImage (
    cl_command_queue command_queue,
    cl_mem src_buffer, cl_mem dst_image,
    size_t src_offset, const size_t dst_origin[3],
    const size_t region[3],
    cl_uint num_events_in_wait_list,
    const cl_event *event_wait_list, cl_event *event)
```

## Map and Unmap Image Objects [5.2.8]

```
void *clEnqueueMapImage (
    cl_command_queue command_queue,
    cl_mem image, cl_bool blocking_map,
    cl_map_flags map_flags, const size_t origin[3],
    const size_t region[3], size_t *image_row_pitch,
    size_t *image_slice_pitch,
    cl_uint num_events_in_wait_list,
    const cl_event *event_wait_list,
    cl_event *event, cl_int *errcode_ret)
```

## Read, Write, Copy Image Objects [5.2.6]

```
cl_int clEnqueueReadImage (
    cl_command_queue command_queue,
    cl_mem image, cl_bool blocking_read,
    const size_t origin[3], const size_t region[3],
    size_t row_pitch, size_t slice_pitch, void *ptr,
    cl_uint num_events_in_wait_list,
    const cl_event *event_wait_list, cl_event *event)

cl_int clEnqueueWriteImage (
    cl_command_queue command_queue,
    cl_mem image, cl_bool blocking_write,
    const size_t origin[3], const size_t region[3],
    size_t input_row_pitch, size_t input_slice_pitch,
    const void *ptr, cl_uint num_events_in_wait_list,
    const cl_event *event_wait_list, cl_event *event)

cl_int clEnqueueCopyImage (
    cl_command_queue command_queue,
    cl_mem src_image, cl_mem dst_image,
    const size_t src_origin[3], const size_t dst_origin[3],
    const size_t region[3],
    cl_uint num_events_in_wait_list,
    const cl_event *event_wait_list, cl_event *event)
```

## Query Image Objects [5.2.9]

```
cl_int clGetMemObjectInfo (cl_mem memobj,
    cl_mem_info param_name,
    size_t param_value_size, void *param_value,
    size_t *param_value_size_ret)

param_name: CL_MEM_TYPE,
CL_MEM_FLAGS,           CL_MEM_SIZE,
CL_MEM_HOST_PTR,        CL_MEM_MAP_COUNT,
CL_MEM_REFERENCE_COUNT, CL_MEM_CONTEXT

cl_int clGetImageInfo (cl_mem image,
    cl_image_info param_name,
    size_t param_value_size, void *param_value,
    size_t *param_value_size_ret)

param_name: CL_IMAGE_FORMAT,
CL_IMAGE_ELEMENT_SIZE,   CL_IMAGE_ROW_PITCH,
CL_IMAGE_SLICE_PITCH,     CL_IMAGE_HEIGHT,
CL_IMAGE_WIDTH,          CL_IMAGE_DEPTH
```

## Image Formats [5.2.4.1, 9.8]

Supported image formats: image\_channel\_order with image\_channel\_data\_type.

### Built-in support:

<b>CL_RGBA</b> : CL_HALF_FLOAT, CL_FLOAT, CL_UNORM_INT{8 16}, CL_SIGNED_INT{8 16 32}, CL_UNSIGNED_INT{8 16 32}
--

<b>CL_BGRA</b> : CL_UNORM_INT8
--------------------------------

### Optional support:

<b>CL_R, CL_A</b> : CL_HALF_FLOAT, CL_FLOAT, CL_UNORM_INT{8 16}, CL_SIGNED_INT{8 16 32}, CL_UNSIGNED_INT{8 16 32}, CL_SNORM_INT{8 16}
---

<b>CL_INTENSITY</b> : CL_HALF_FLOAT, CL_FLOAT, CL_UNORM_INT{8 16}, CL_SNORM_INT{8 16}
---

<b>CL_LUMINANCE</b> : CL_UNORM_INT{8 16}, CL_HALF_FLOAT, CL_FLOAT, CL_SNORM_INT{8 16}
---

<b>CL_RG, CL_RA</b> : CL_HALF_FLOAT, CL_FLOAT, CL_UNORM_INT{8 16}, CL_SIGNED_INT{8 16 32}, CL_UNSIGNED_INT{8 16 32}, CL_SNORM_INT{8 16}
---

<b>CL_RGB</b> : CL_UNORM_SHORT_{555 565}, CL_UNORM_INT_101010
---

<b>CL_ARGB</b> : CL_UNORM_INT8, CL_SIGNED_INT8, CL_UNSIGNED_INT8, CL_SNORM_INT8
---

<b>CL_BGRA</b> : CL_SIGNED_INT8, CL_UNSIGNED_INT8, CL_SNORM_INT8
--

# Lesen und Schreiben von Bildern

## Image Read and Write Built-in Functions [6.11.8, 9.8]

The built-in functions defined in this section can only be used with image memory objects created with `clCreateImage2D` or `clCreateImage3D`. **OPT** = Optional function.

<code>float4 read_imagef (image2d_t image, sampler_t sampler, int2 coord)</code>		
<code>float4 read_imagef (image2d_t image, sampler_t sampler, float2 coord)</code>		
<code>int4 read_imagei (image2d_t image, sampler_t sampler, int2 coord)</code>		
<code>int4 read_imagei (image2d_t image, sampler_t sampler, float2 coord)</code>		
<code>unsigned int4 read_imageui (image2d_t image, sampler_t sampler, int2 coord)</code>		
<code>unsigned int4 read_imageui (image2d_t image, sampler_t sampler, float2 coord)</code>		
<code>half4 read_imaged (image2d_t image, sampler_t sampler, int2 coord)</code>	<b>OPT</b>	
<code>half4 read_imaged (image2d_t image, sampler_t sampler, float2 coord)</code>	<b>OPT</b>	
<code>void write_imagef (image2d_t image, int2 coord, float4 color)</code>		
<code>void write_imagei (image2d_t image, int2 coord, int4 color)</code>		
<code>void write_imageui (image2d_t image, int2 coord, unsigned int4 color)</code>		
<code>void write_imaged (image2d_t image, int2 coord, half4 color)</code>	<b>OPT</b>	
<code>float4 read_imaged (image3d_t image, sampler_t sampler, int4 coord)</code>		
<code>float4 read_imaged (image3d_t image, sampler_t sampler, float4 coord)</code>		
<code>int4 read_imagei (image3d_t image, sampler_t sampler, int4 coord)</code>		
<code>int4 read_imagei (image3d_t image, sampler_t sampler, float4 coord)</code>		
<code>unsigned int4 read_imageui (image3d_t image, sampler_t sampler, int4 coord)</code>		
<code>unsigned int4 read_imageui (image3d_t image, sampler_t sampler, float4 coord)</code>		
<code>half4 read_imaged (image3d_t image, sampler_t sampler, int4 coord)</code>	<b>OPT</b>	
<code>half4 read_imaged (image3d_t image, sampler_t sampler, float4 coord)</code>	<b>OPT</b>	
<code>int get_image_width (image2d_t image)</code>		2D or 3D image width in pixels
<code>int get_image_width (image3d_t image)</code>		
<code>int get_image_height (image2d_t image)</code>		2D or 3D image height in pixels
<code>int get_image_height (image3d_t image)</code>		
<code>int get_image_depth (image3d_t image)</code>		3D image depth in pixels
<code>int get_image_channel_data_type (image2d_t image)</code>		image channel data type
<code>int get_image_channel_data_type (image3d_t image)</code>		
<code>int get_image_channel_order (image2d_t image)</code>		image channel order
<code>int get_image_channel_order (image3d_t image)</code>		
<code>int2 get_image_dim (image2d_t image)</code>		2D image width and height
<code>int4 get_image_dim (image3d_t image)</code>		3D image width, height, and depth
<code>void write_imaged (image3d_t image, int4 coord, half4 color)</code>	<b>OPT</b>	Writes <code>color</code> value to $(x, y, z)$ location specified by <code>coord</code> in the 3D image.
<code>void write_imaged (image3d_t image, int4 coord, float4 color)</code>	<b>OPT</b>	Writes <code>color</code> at <code>coord</code> in the 3D image.
<code>void write_imagei (image3d_t image, int4 coord, int4 color)</code>	<b>OPT</b>	Include this pragma to enable these functions:
<code>void write_imageui (image3d_t image, int4 coord, unsigned int4 color)</code>	<b>OPT</b>	#pragma OPENCL EXTENSION cl_khr_3d_image_writes : enable

# Schnittstelle zu OpenGL

## OpenCL/OpenGL Sharing APIs [Appendix B]

Creating OpenCL memory objects from OpenGL objects using the functions `clCreateFromGLBuffer`, `clCreateFromGLTexture2D`, `clCreateFromGLTexture3D`, or `clCreateFromGLRenderbuffer` ensures that the underlying storage of that OpenGL object will not be deleted while the corresponding OpenCL memory object still exists. (Items shown in red are optional)

### CL Buffer Objects > GL Buffer Objects [B.1.1]

```
cl_mem clCreateFromGLBuffer (cl_context context,
    cl_mem_flags flags, GLuint bufobj,
    int *errcode_ret)
flags: CL_MEM_READ_ONLY, CL_MEM_WRITE_ONLY,
CL_MEM_READ_WRITE
```

### CL Image Objects > GL Textures [B.1.2]

```
cl_mem clCreateFromGLTexture2D (
    cl_context context, cl_mem_flags flags,
    GLenum target, GLint mipmap, GLuint texture,
    int *errcode_ret)
```

*flags:* (Same as for `clCreateFromGLBuffer`)

```
target: GL_TEXTURE_2D, GL_TEXTURE_RECTANGLE,
GL_TEXTURE_CUBE_MAP_POSITIVE_X | Y | Z),
GL_TEXTURE_CUBE_MAP_NEGATIVE_X | Y | Z)
```

```
cl_mem clCreateFromGLTexture3D (
    cl_context context, cl_mem_flags flags,
    GLenum target, GLint mipmap, GLuint texture,
    int *errcode_ret)
```

*flags:* (Same as for `clCreateFromGLBuffer`)

*target:* GL\_TEXTURE\_3D

### CL Image Objects > GL Renderbuffers [B.1.3]

```
cl_mem clCreateFromGLRenderbuffer (
    cl_context context, cl_mem_flags flags,
    GLuint renderbuffer, int *errcode_ret)
```

*flags:* (Same as for `clCreateFromGLBuffer`)

### Query Information [B.1.4]

```
cl_int clGetGLObjectInfo (cl_mem memobj,
    cl_object_type *gl_object_type,
    GLuint *gl_object_name)
```

```
gl_object_type: CL_GL_OBJECT_BUFFER,
CL_GL_OBJECT_TEXTURE2D,
CL_GL_OBJECT_TEXTURE_RECTANGLE,
CL_GL_OBJECT_TEXTURE3D,
CL_GL_OBJECT_RENDERBUFFER
```

```
cl_int clGetGLTextureInfo (cl_mem memobj,
    cl_gl_texture_info param_name,
    size_t param_value_size, void *param_value,
    size_t *param_value_size_ret)
param_name: CL_GL_TEXTURE_TARGET,
CL_GL_MIPMAP_LEVEL
```

### Share Objects [B.1.5]

```
cl_int clEnqueueAcquireGLObjects (
    cl_command_queue command_queue,
    cl_uint num_objects, const cl_mem *mem_objects,
    cl_uint num_events_in_wait_list,
    const cl_event *event_wait_list, cl_event *event)
```

```
cl_int clEnqueueReleaseGLObjects (
    cl_command_queue command_queue,
    cl_uint num_objects,
    const cl_mem *mem_objects,
    cl_uint num_events_in_wait_list,
    const cl_event *event_wait_list, cl_event *event)
```

### Querying CL Devices in GL Context [9.11]

```
cl_int clGetGLContextInfoKHR (
    const cl_context_properties *properties,
    cl_gl_context_info param_name,
    size_t param_value_size, void *param_value,
    size_t *param_value_size_ret)
```

*param\_name:* CL\_DEVICES\_FOR\_GL\_CONTEXT\_KHR,  
CL\_CURRENT\_DEVICE\_FOR\_GL\_CONTEXT\_KHR

# OpenCL Empfehlungen

- Gute Parallelisierung eines sequentiellen Algorithmus  
**(Datenparallelität mit ‚leichtem‘ Kernel ist optimal)**
- **Datentransfer** zwischen Hauptcomputer und Gerät **minimieren**
- Verwendung von globalem Speicher minimieren.  
**Privaten und lokalen Speicher bevorzugen**
- **Verschiedene Ausführungspfade** im Kernel **vermeiden**
- Verwenden der **optimierten OpenCL Mathe-Bibliothek**  
(evtl. auch mit reduzierter Genauigkeit)

**(Optimierung:** NVIDIA OpenCL Best Practices Guide 1.0, PDF)